

NoSQL vs. NewSQL

Demystifying the Zoo of Contemporary Database Systems

ComputeFest

Niv Dayan

12 January, 2017



NoSQL

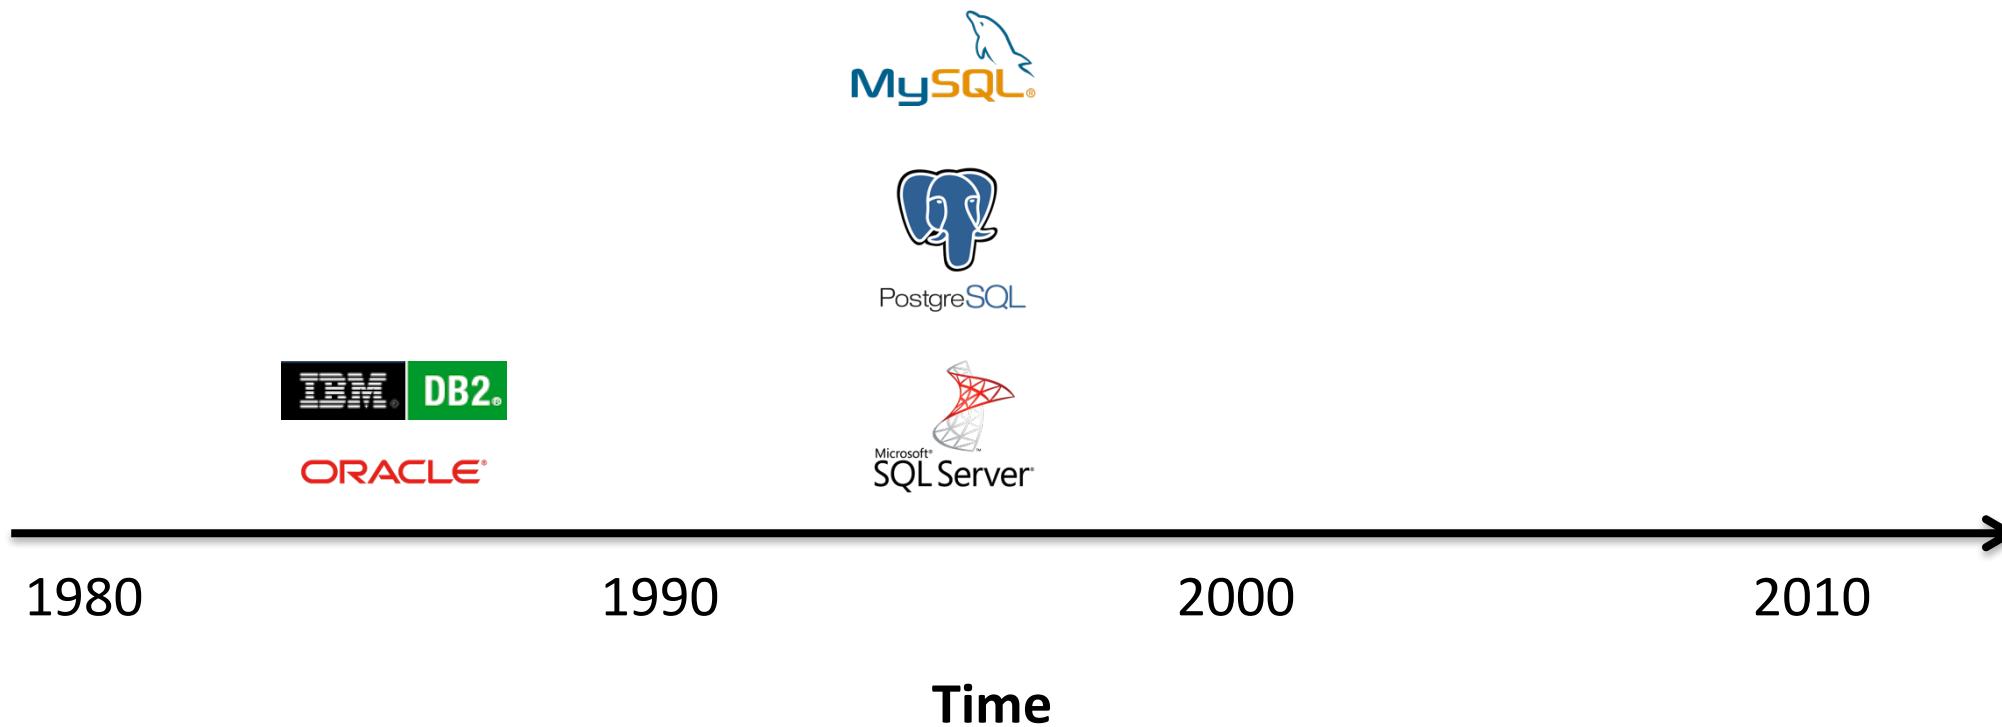
NewSQL



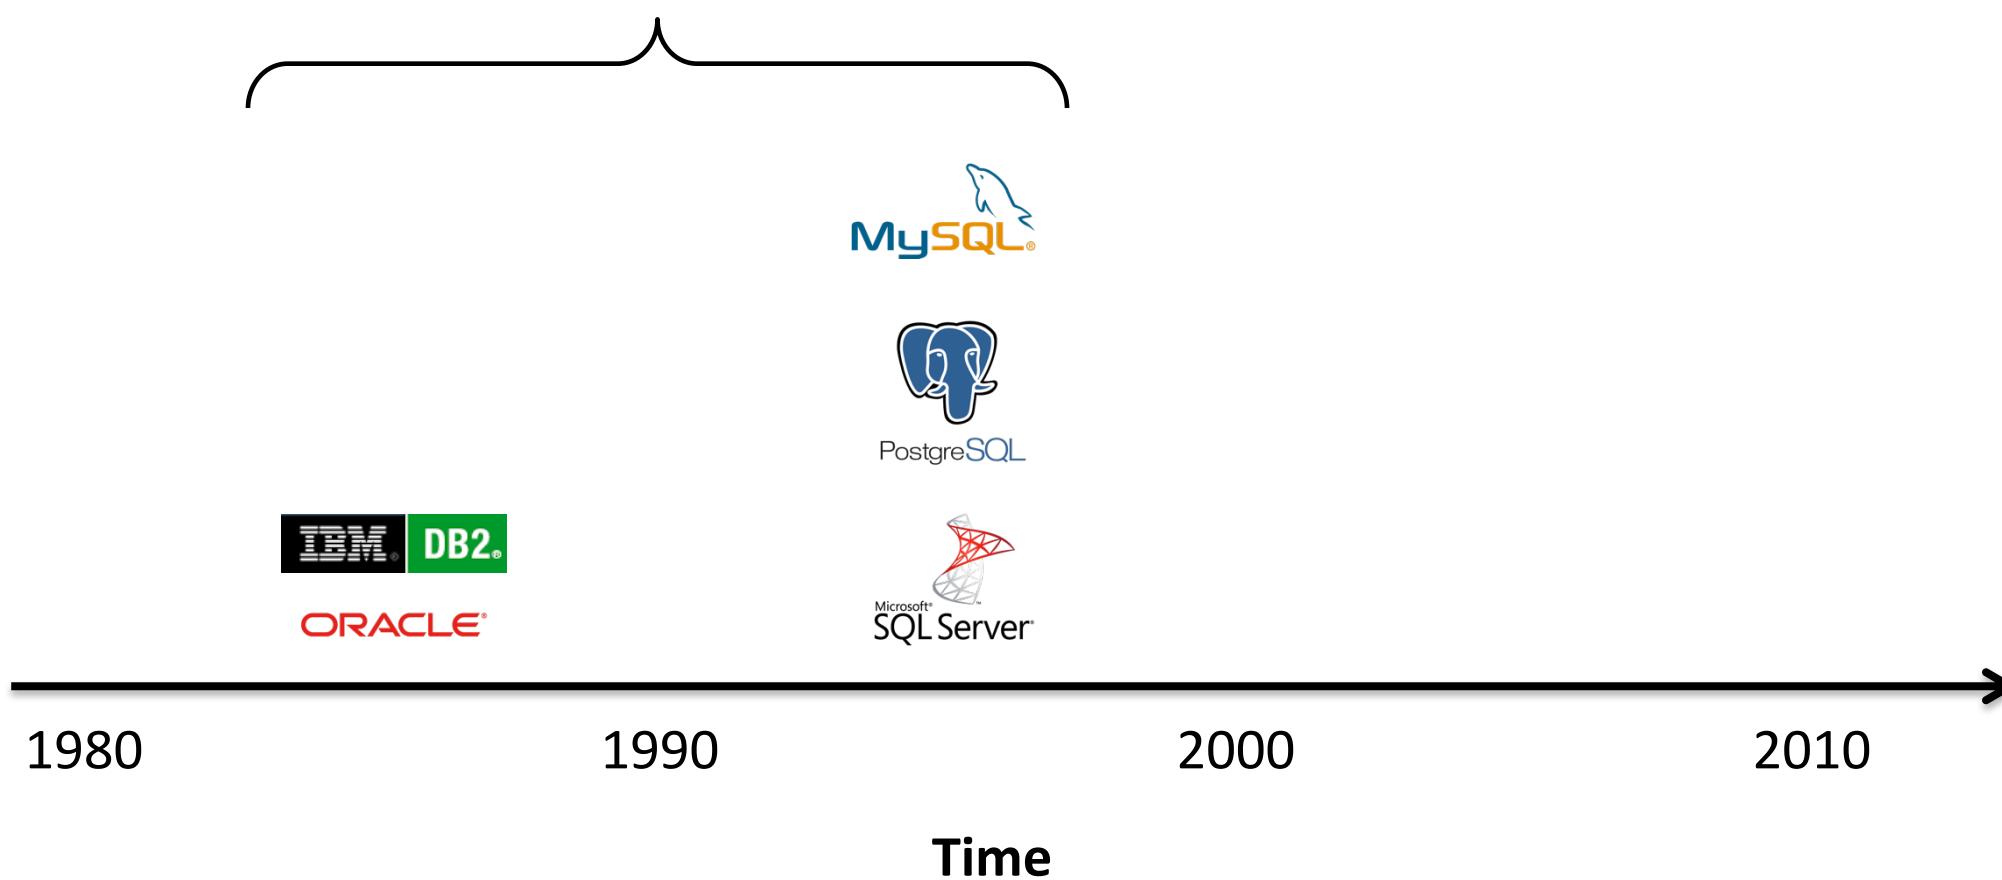
NoSQL

NewSQL

Buzzwords => Principles

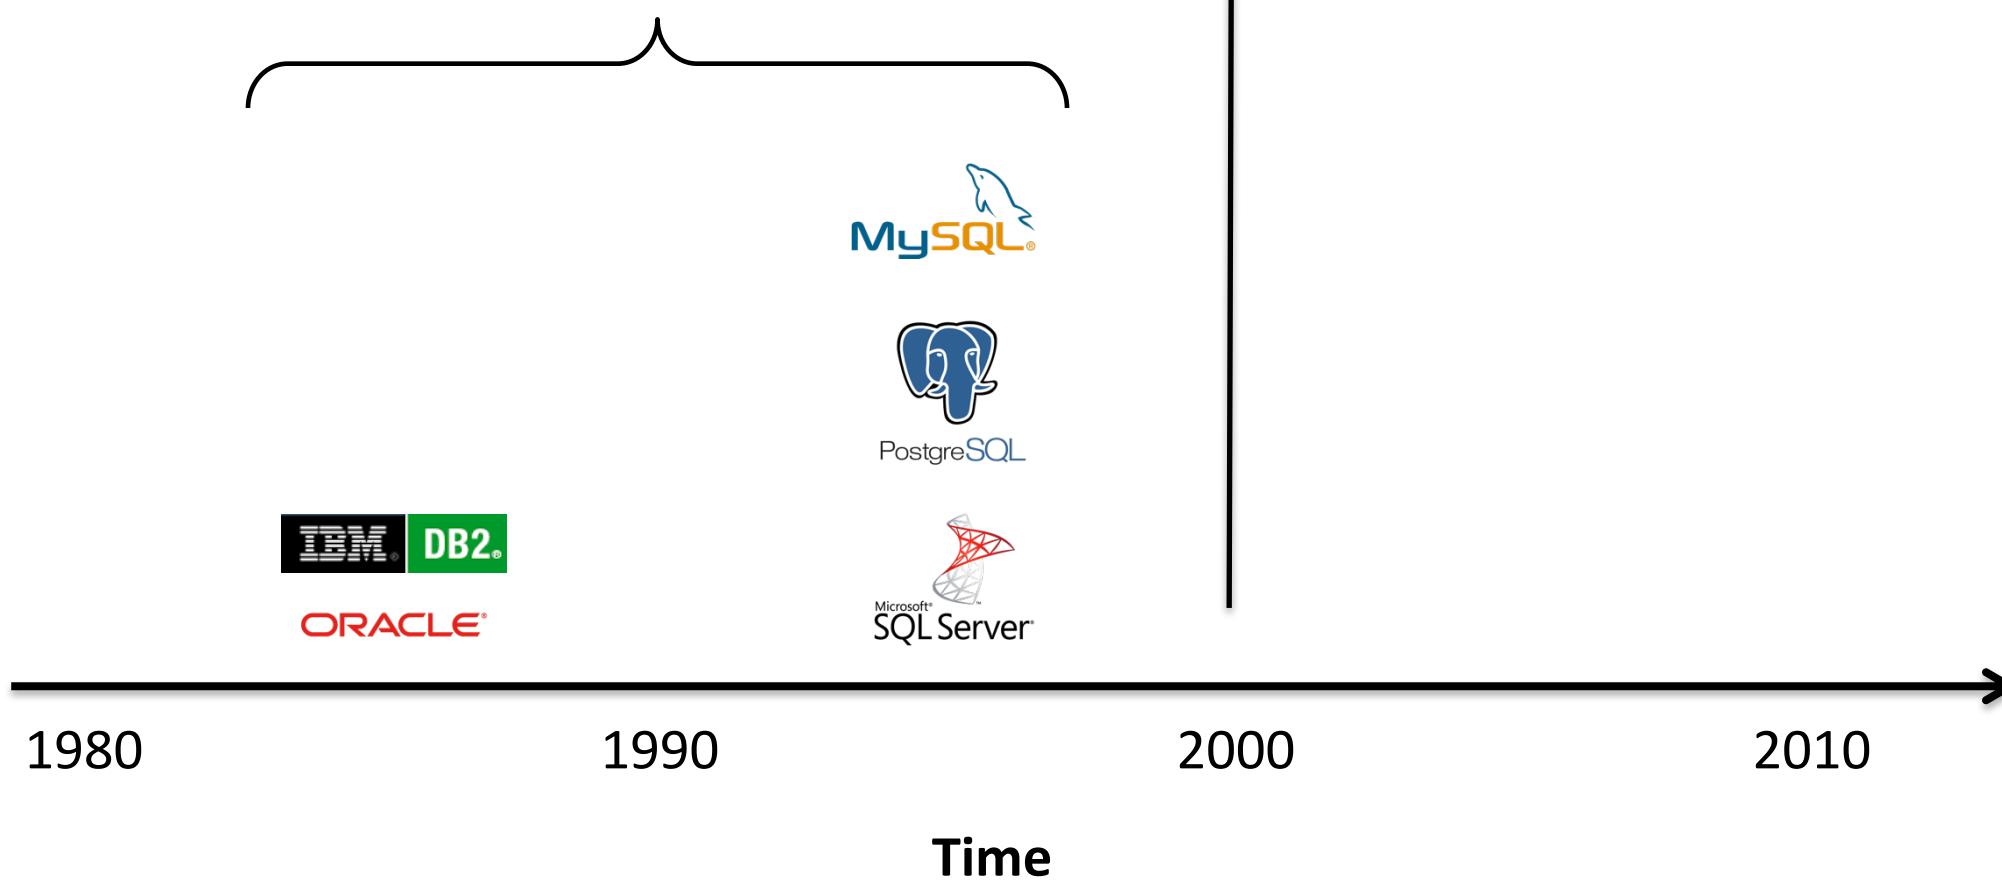


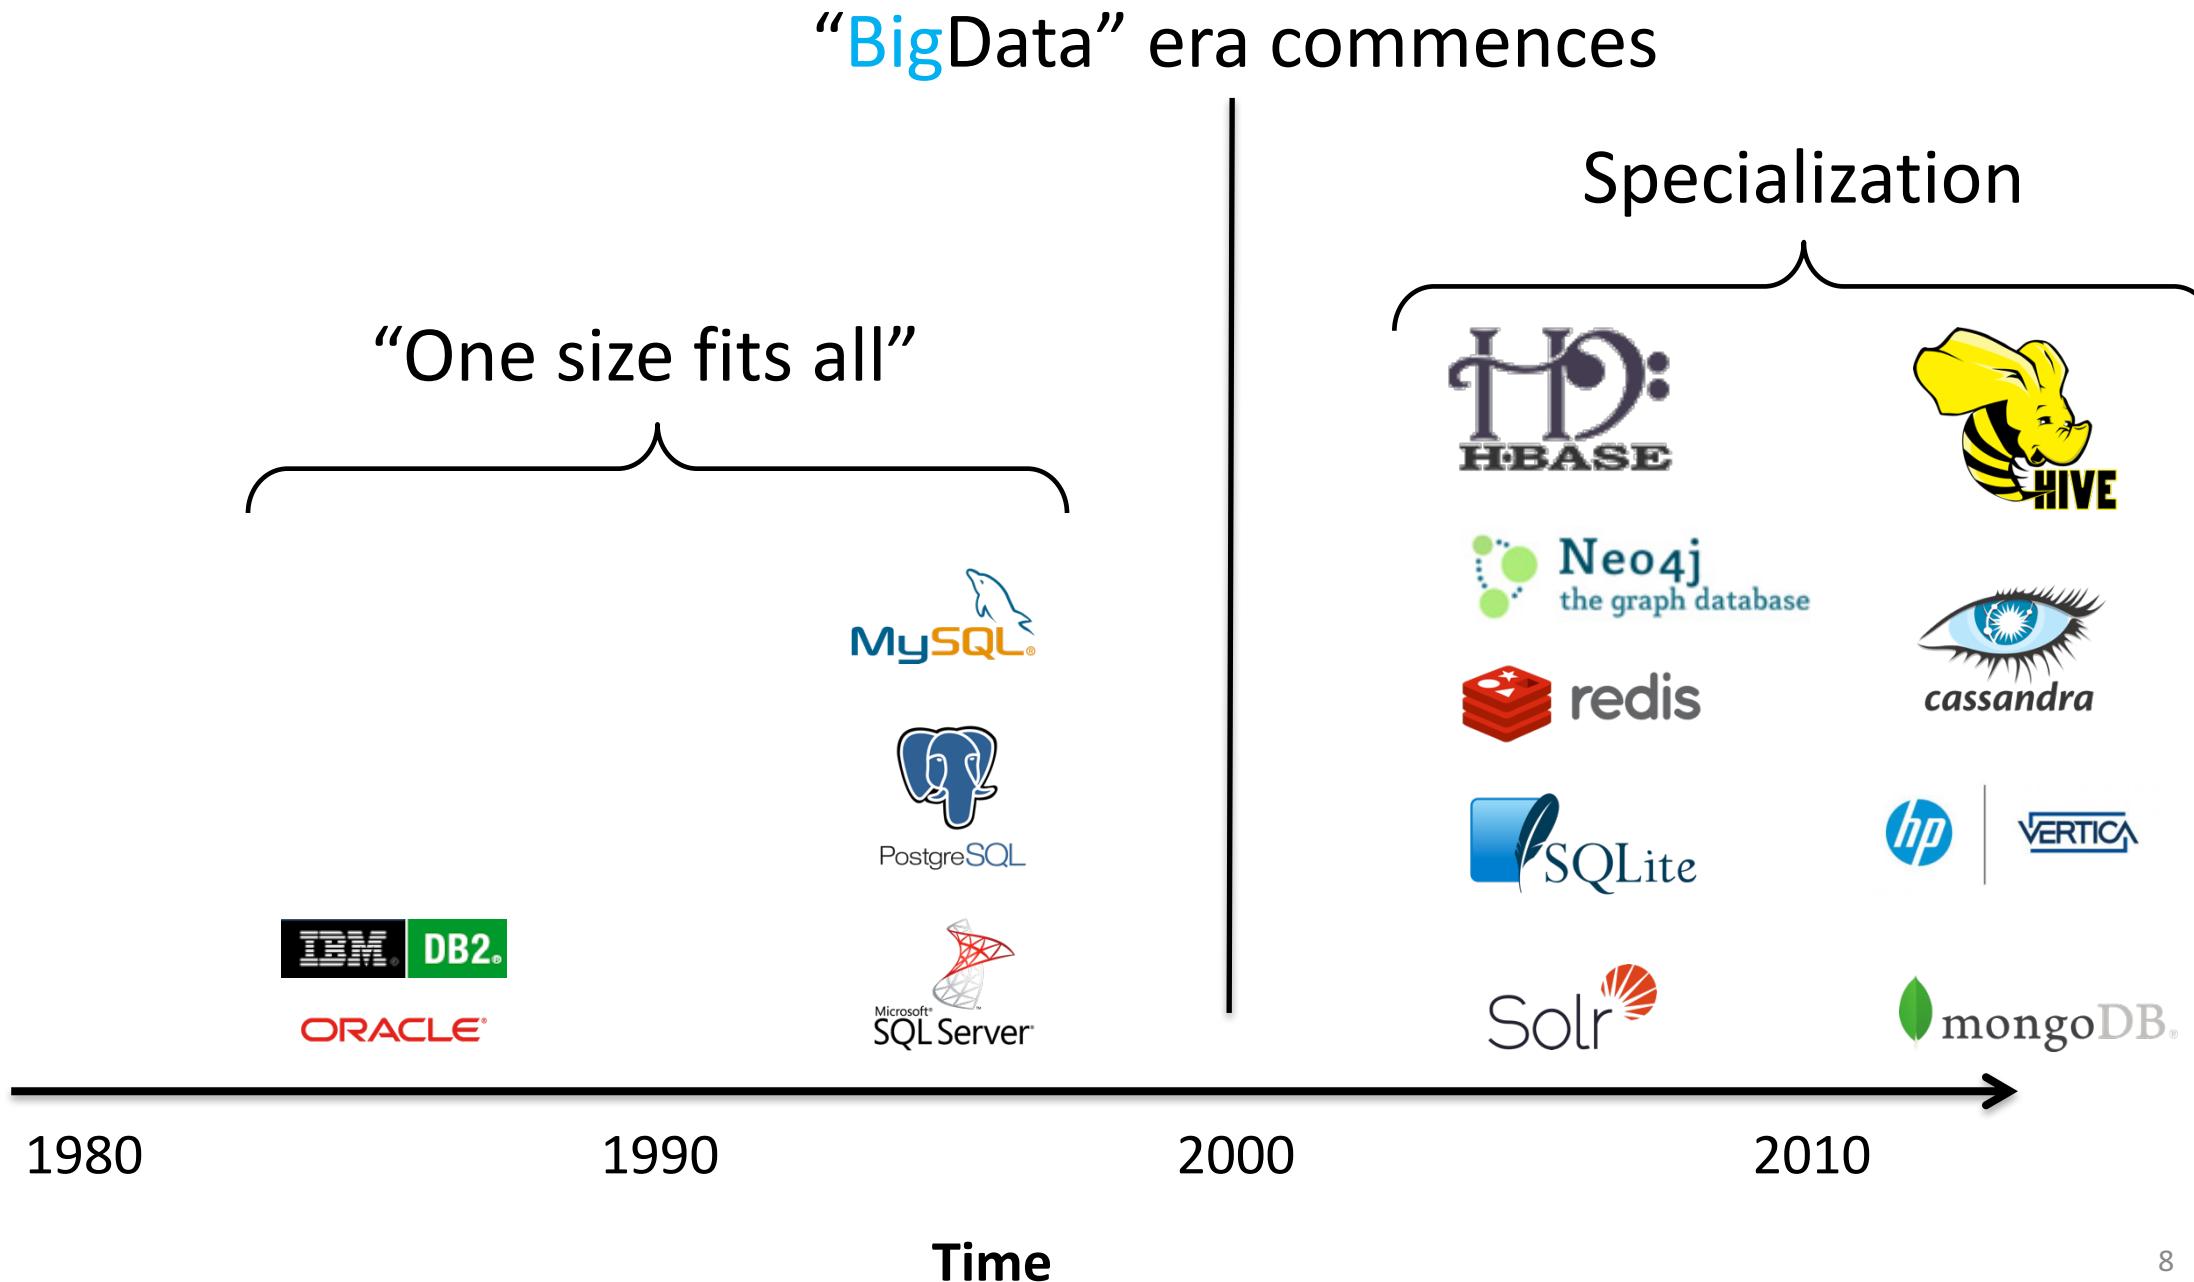
“One size fits all”



“BigData” era commences

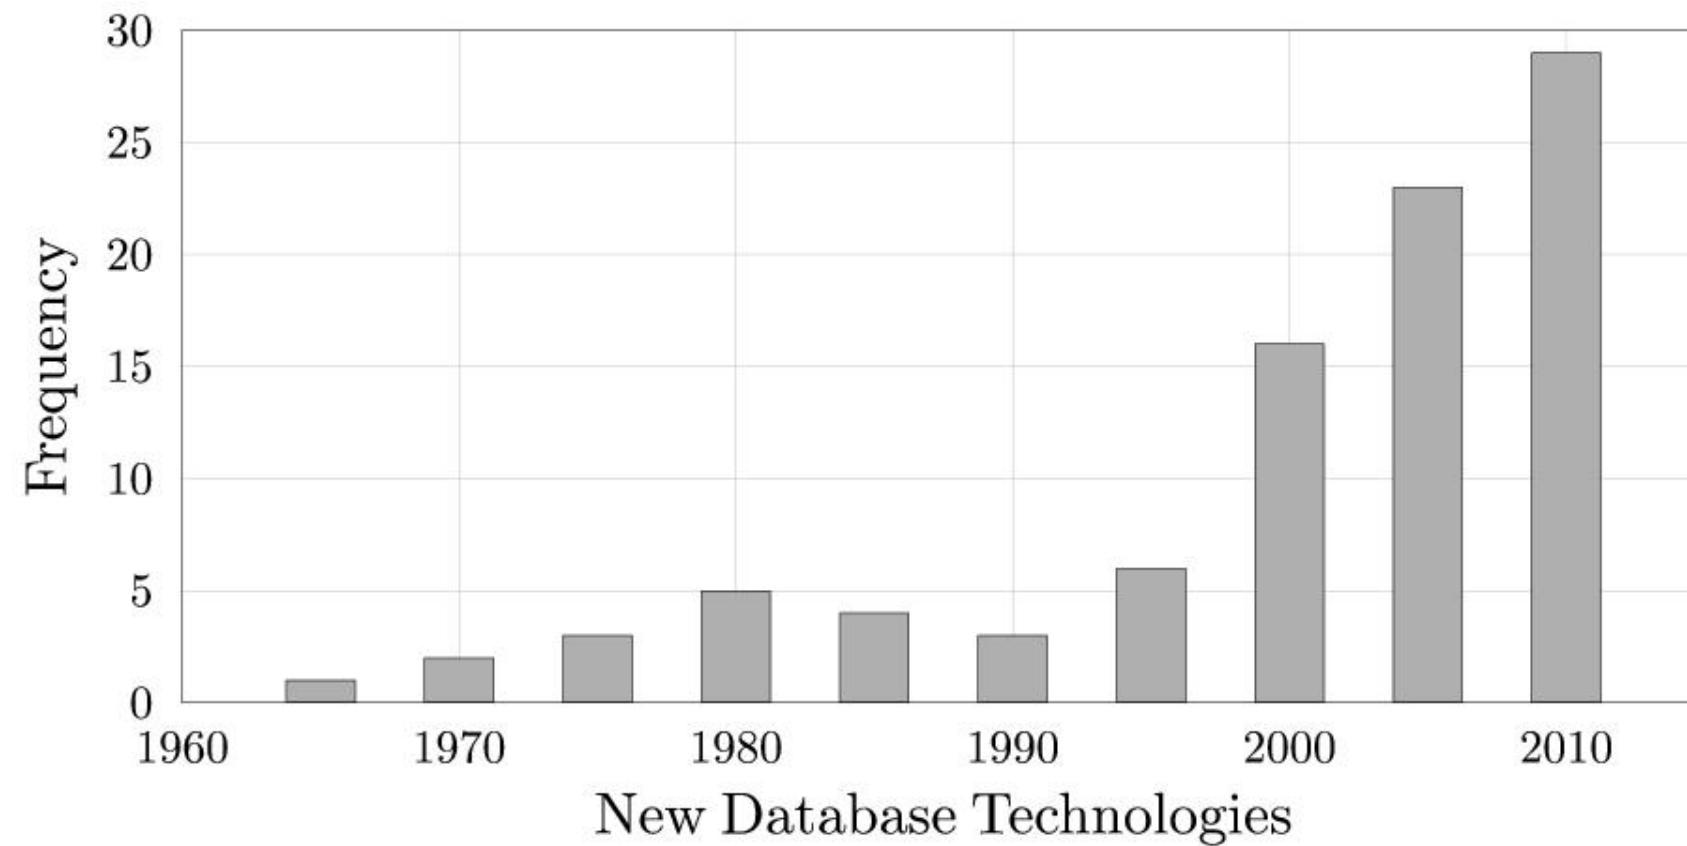
“One size fits all”

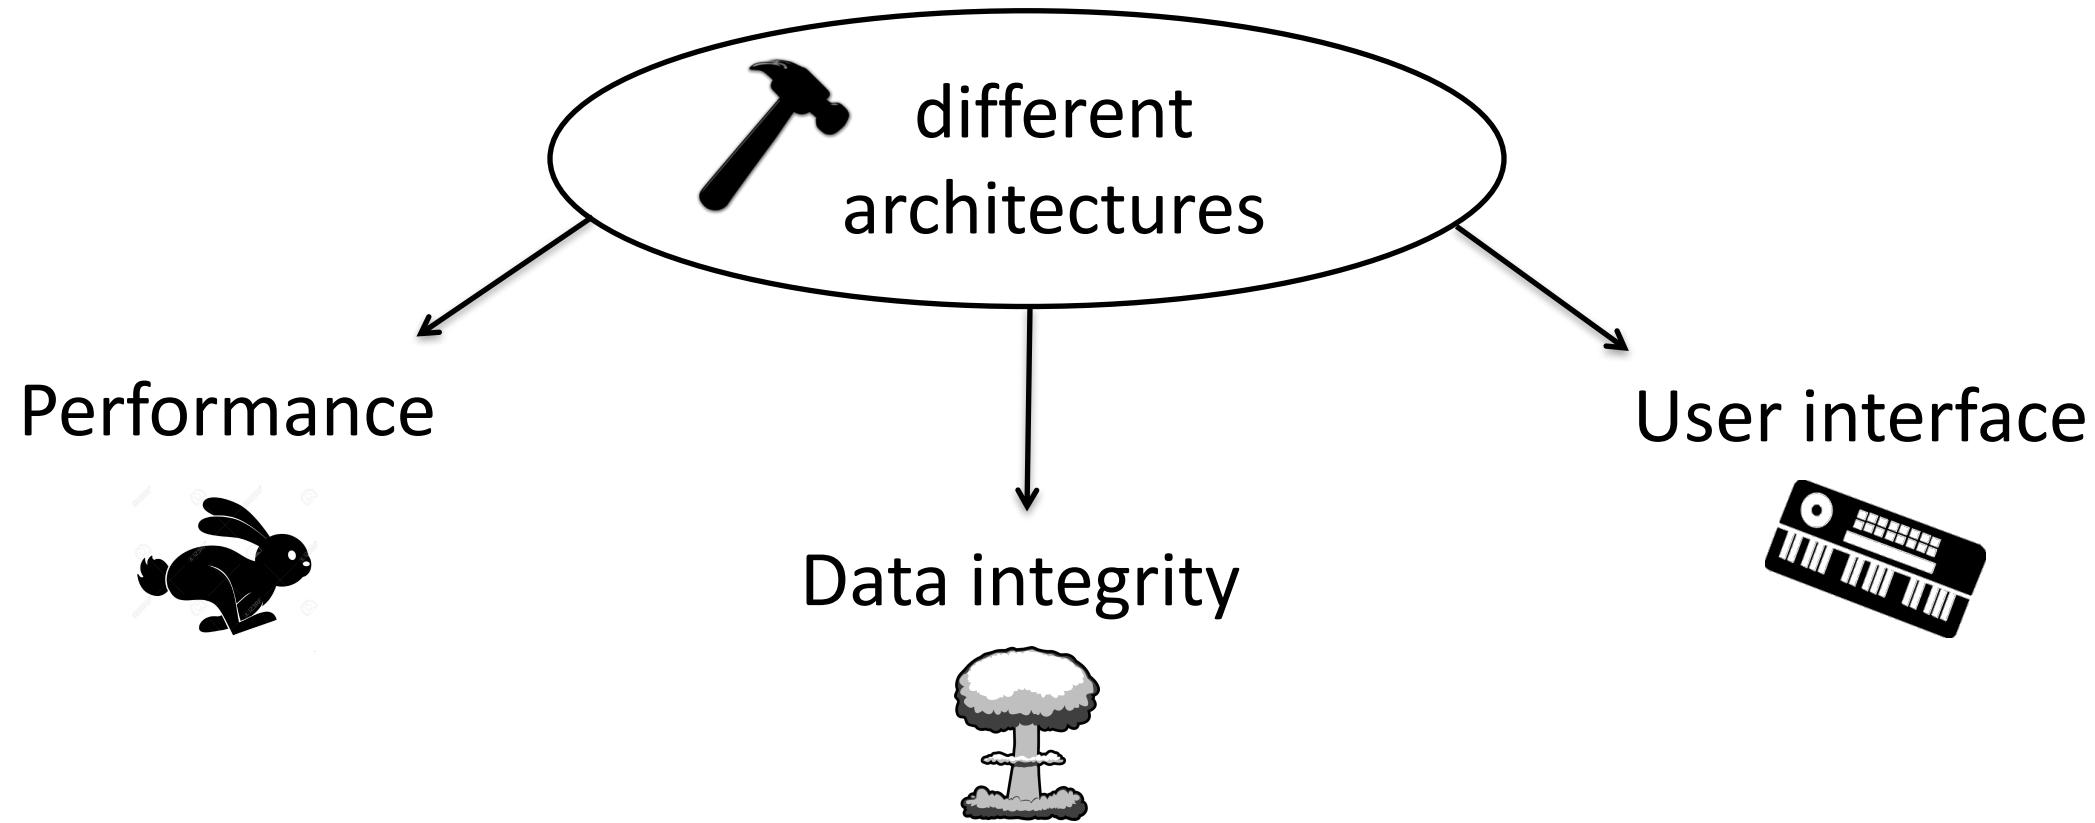




Most database systems were created recently!

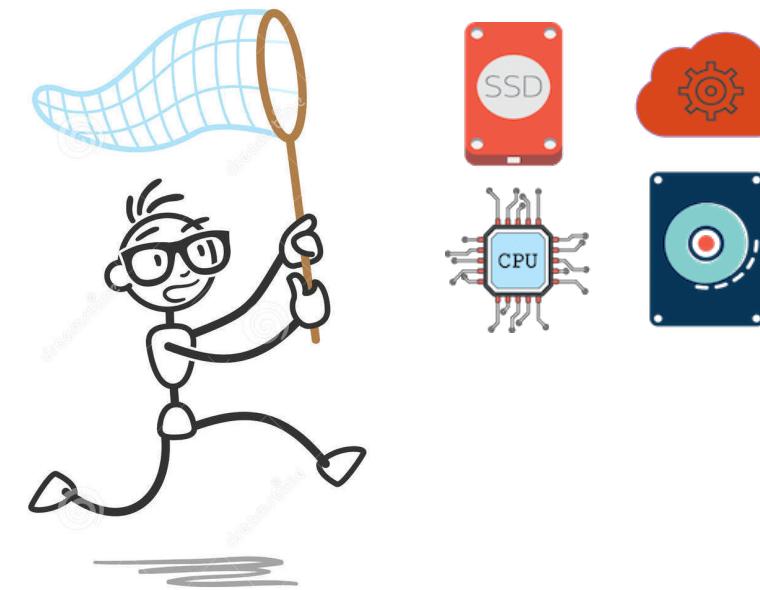
(db-engines.com)





“Which database system is right for me?”

Theme: databases trends follow hardware trends



Claim: modern databases are adaptations

Different backgrounds, No Assumptions



Agenda

1. Traditional databases
2. BigData challenges
3. Analytical databases
4. NoSQL databases
5. NewSQL databases



Total time: 90 minutes

Why MongoDB?

Rank			DBMS
Jan 2017	Dec 2016	Jan 2016	
1.	1.	1.	Oracle 
2.	2.	2.	MySQL 
3.	3.	3.	Microsoft SQL Server
4.	↑ 5.	4.	MongoDB 
5.	↓ 4.	5.	PostgreSQL
6.	6.	6.	DB2
7.	7.	↑ 8.	Cassandra 
8.	8.	↓ 7.	Microsoft Access
9.	9.	↑ 10.	Redis 
10.	10.	↓ 9.	SQLite

(db-engines.com)

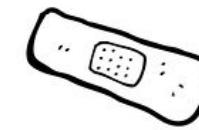
Traditional Databases

Database Design Goals

1. Speed
2. Affordability
3. Resilience to system failure

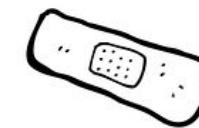


\$



Database Design Goals

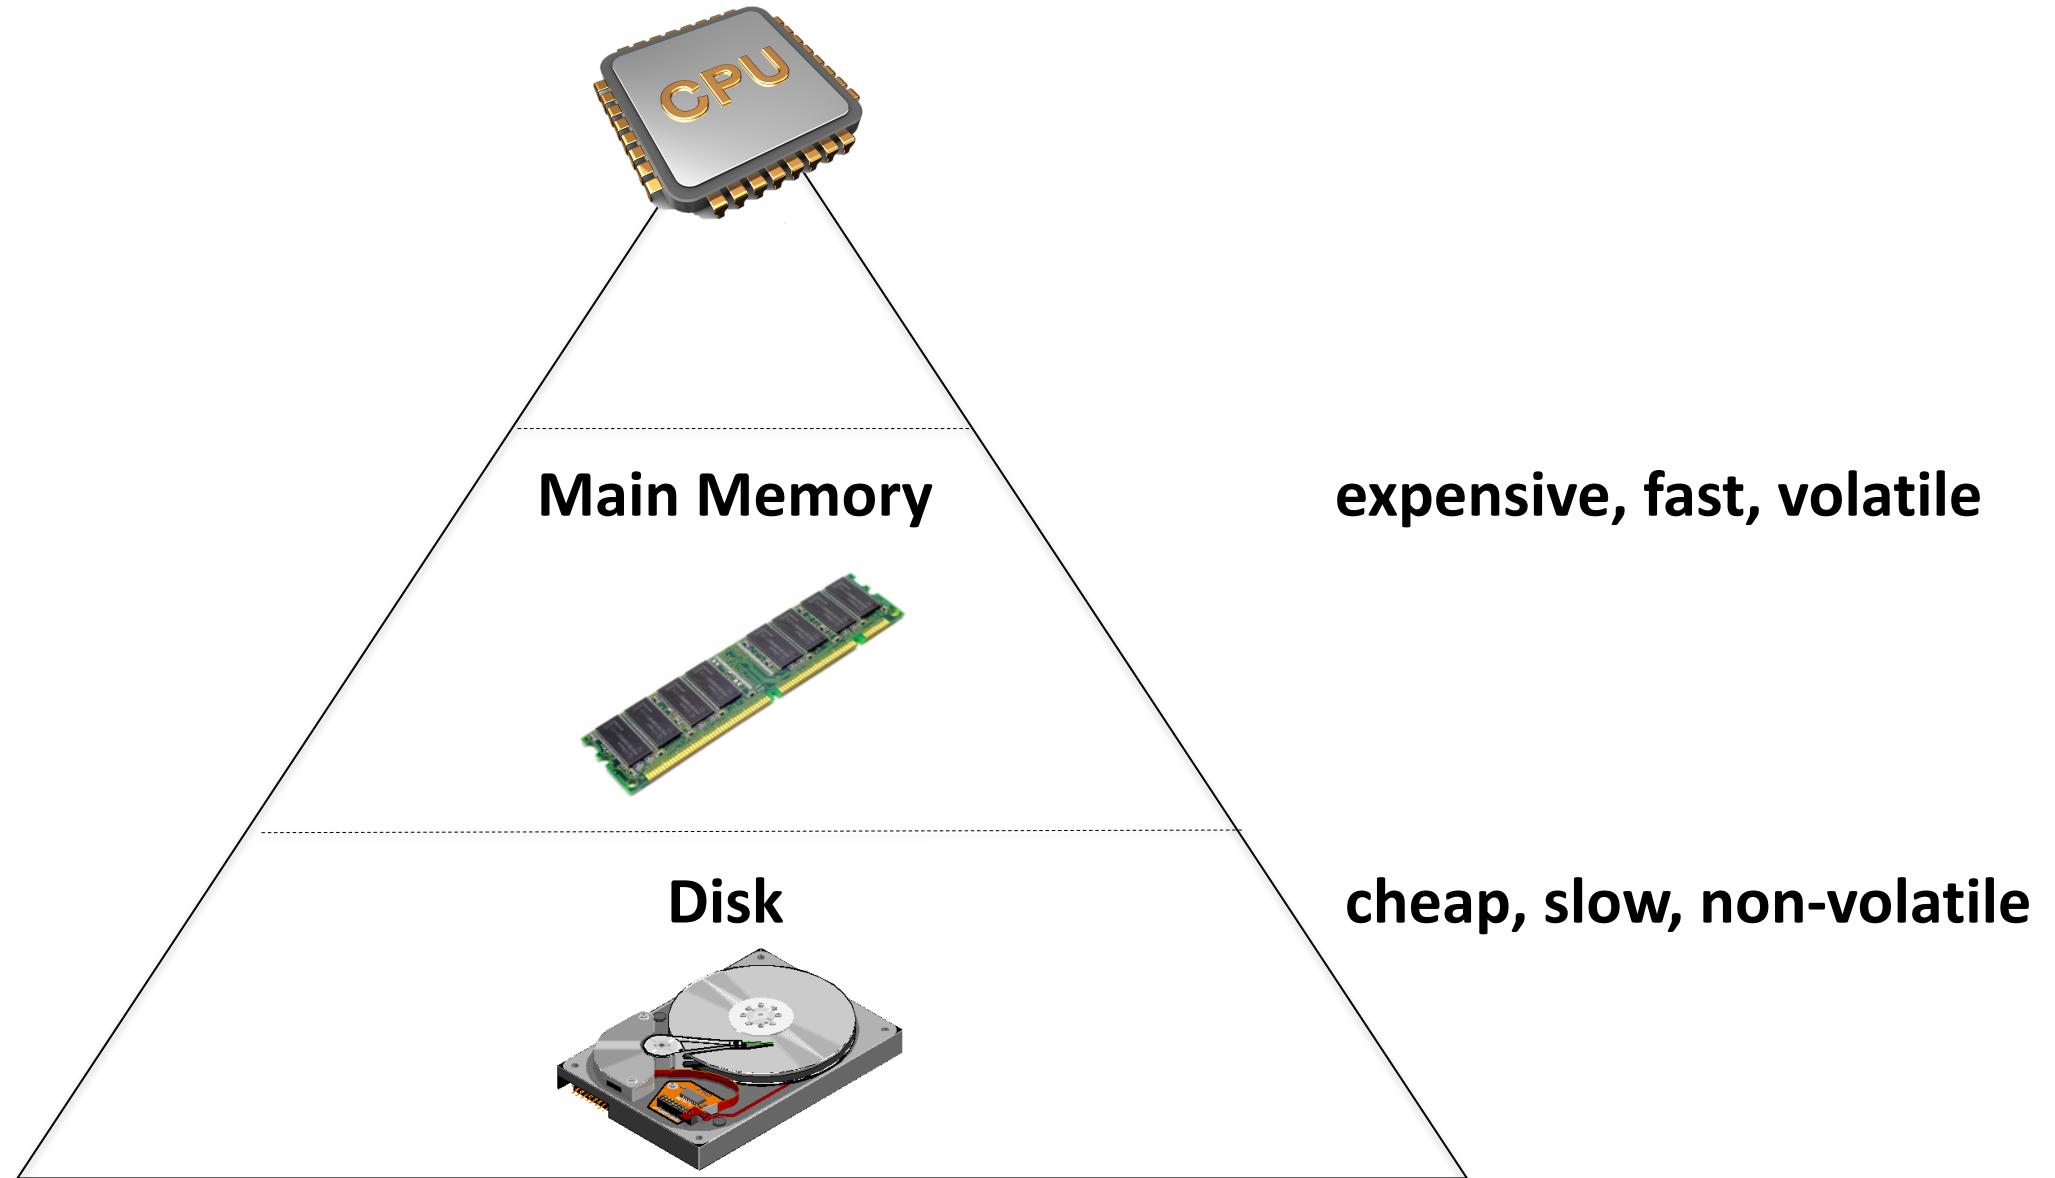
1. Speed
2. Affordability
3. Resilience to system failure



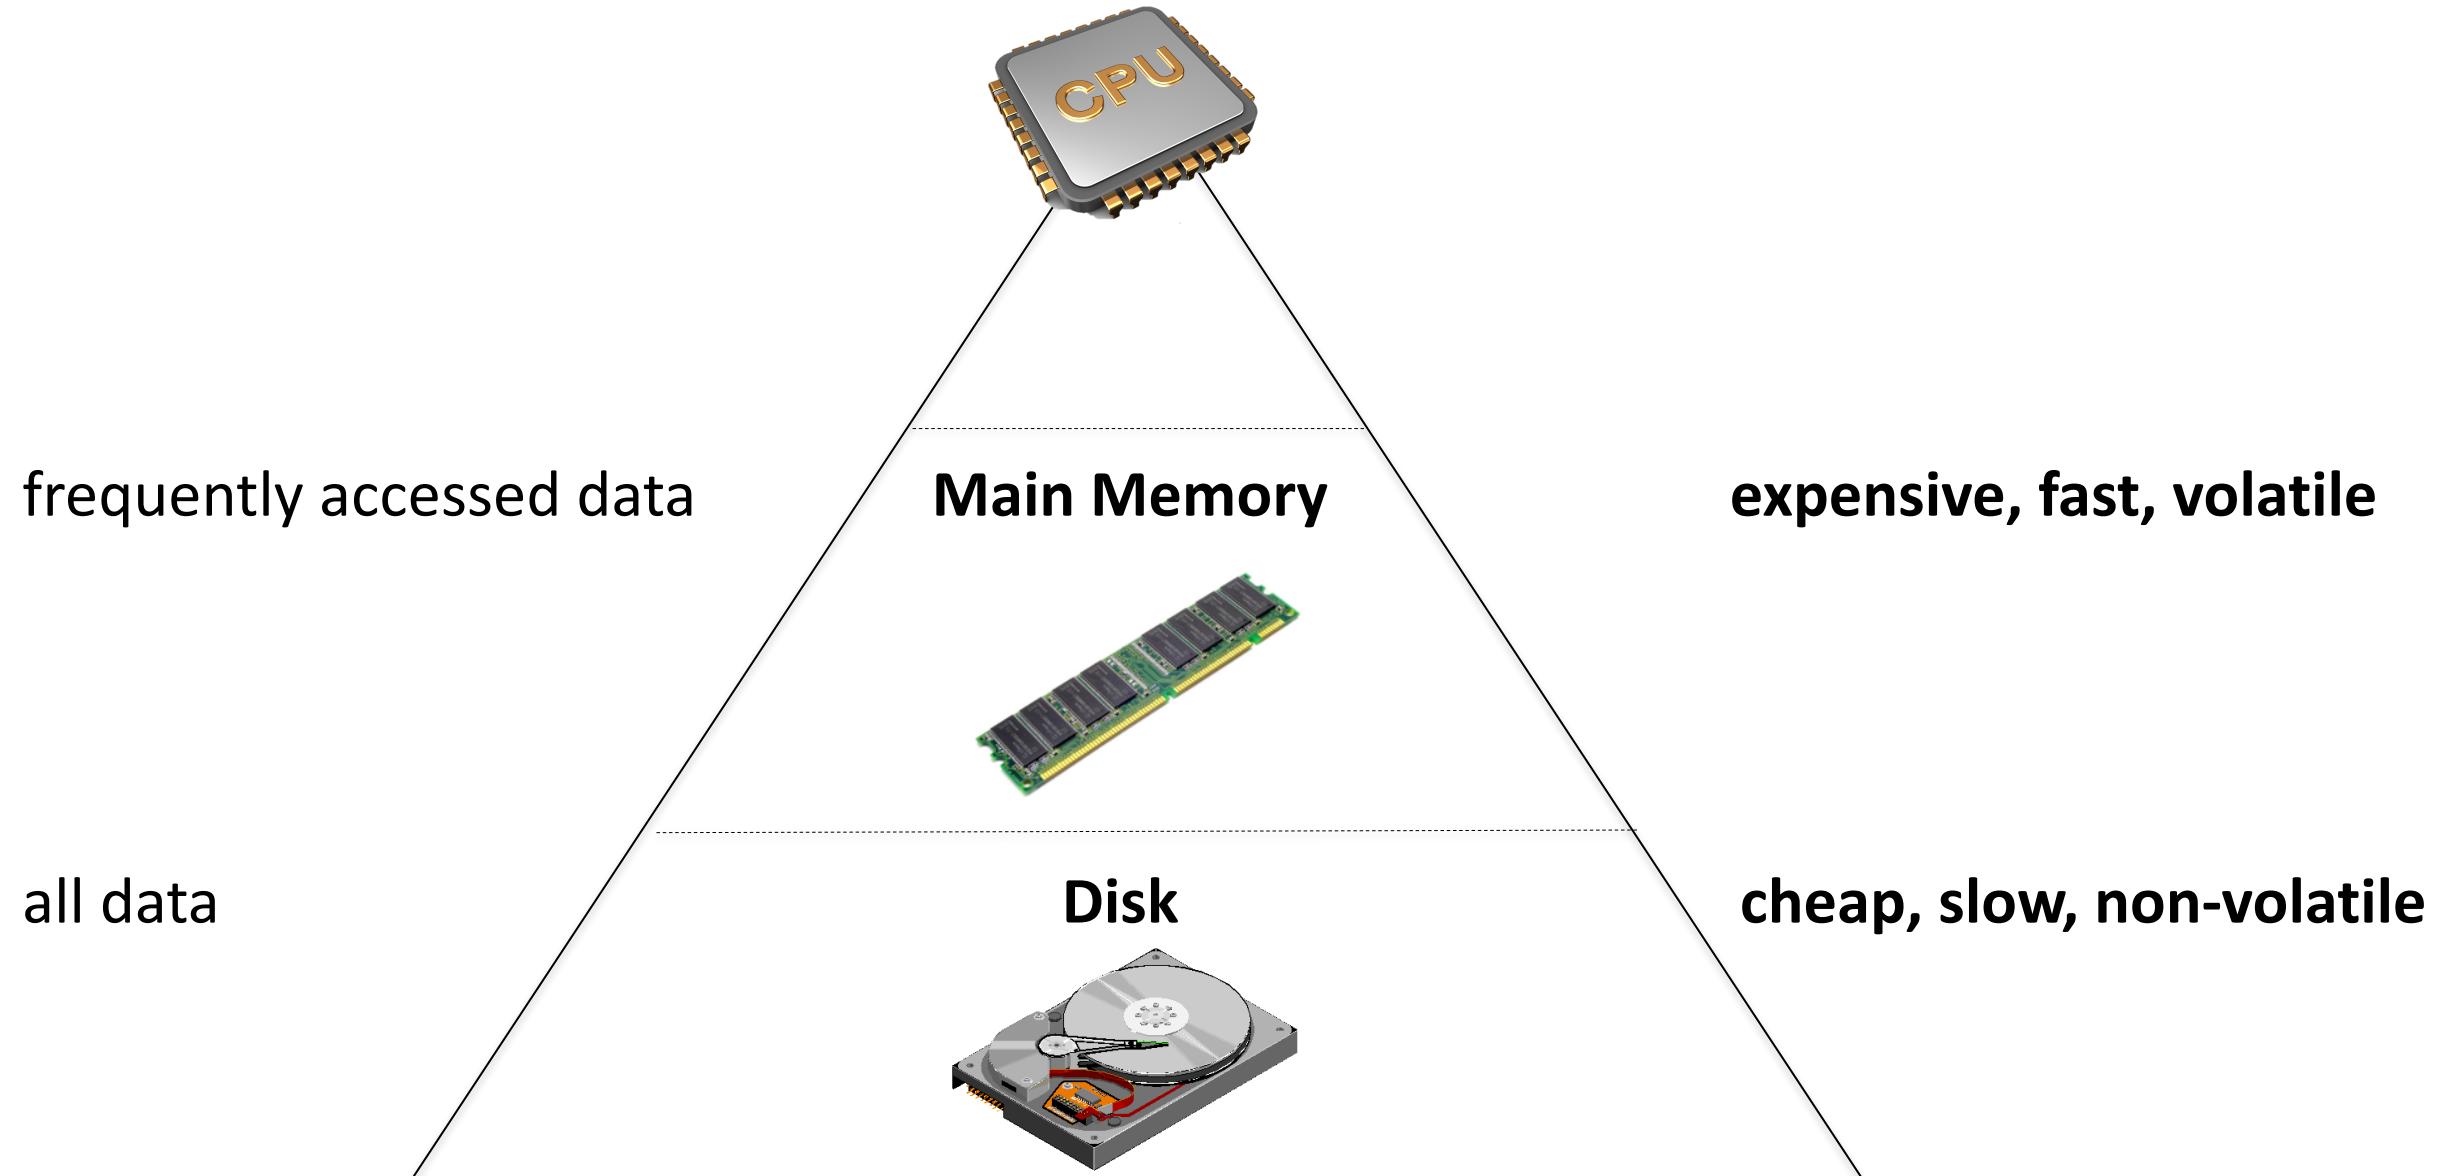
How you achieve them depends on storage hardware



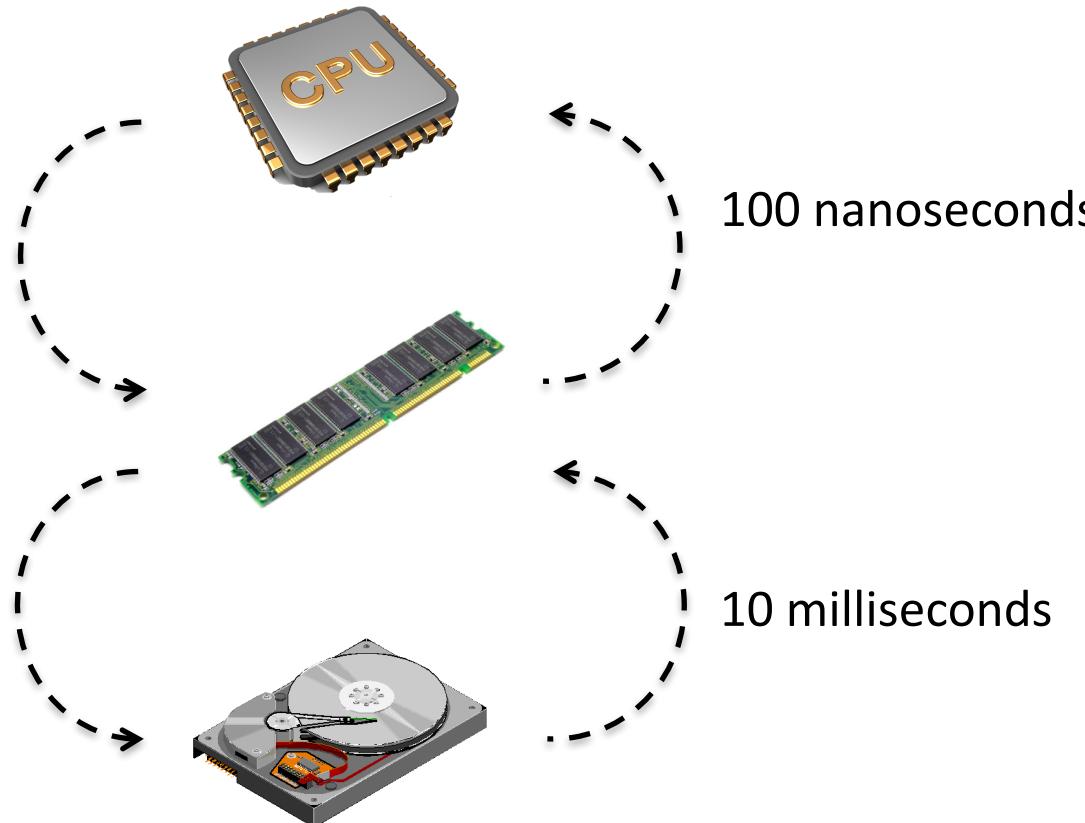
The Memory Hierarchy



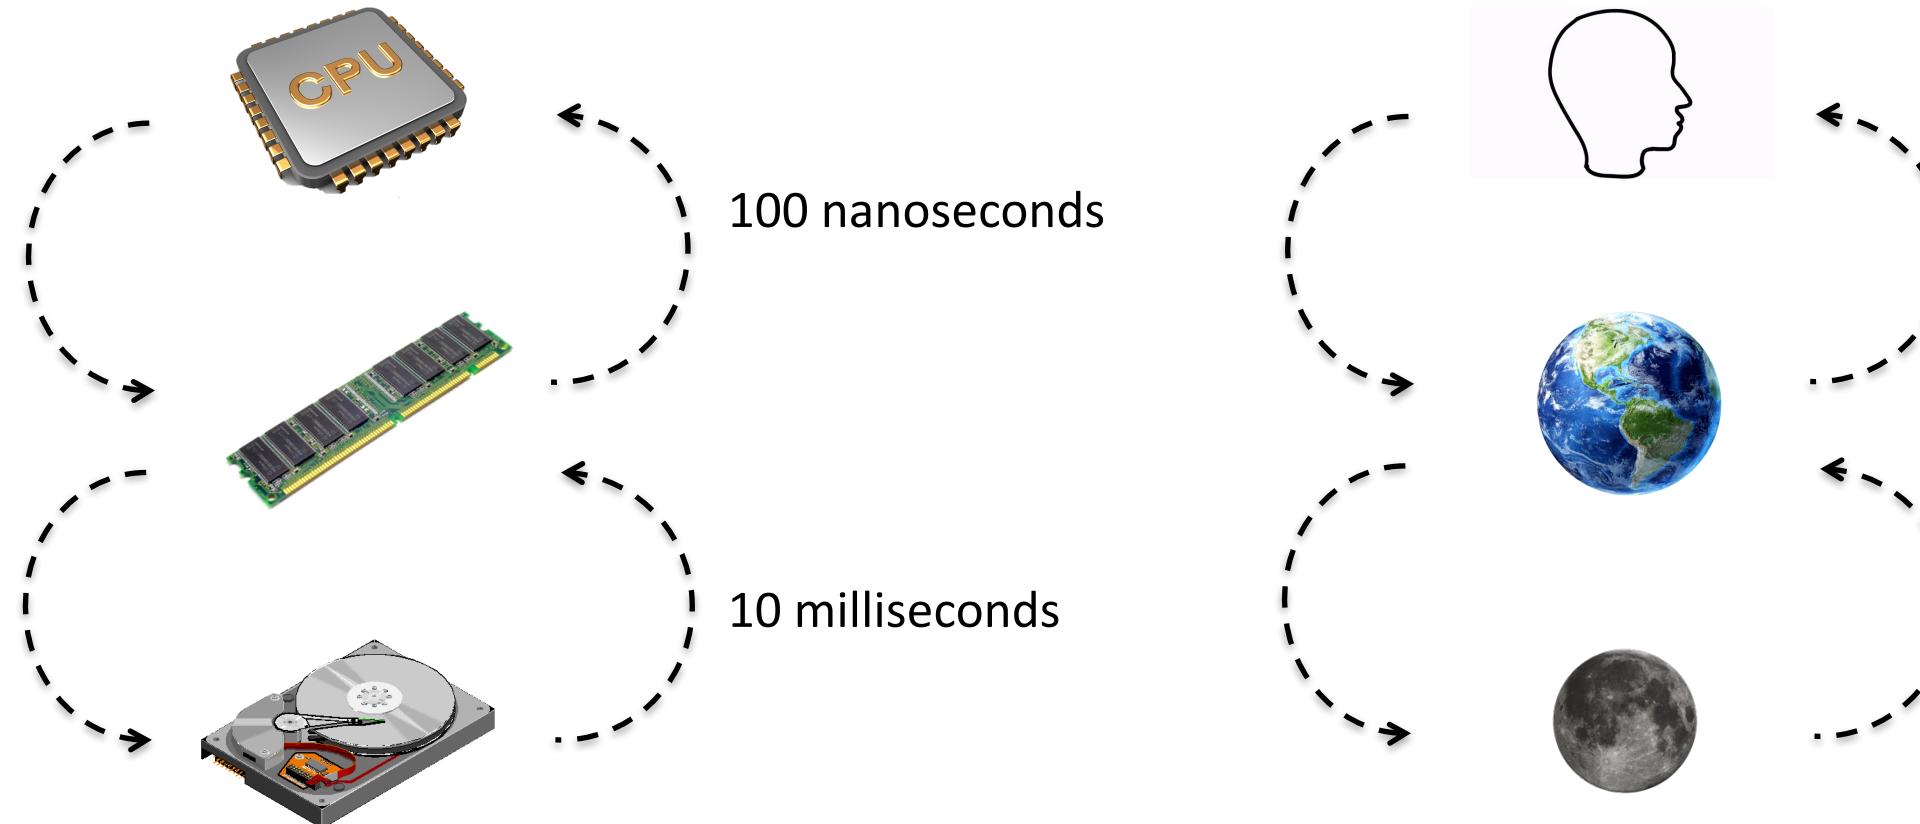
The Memory Hierarchy



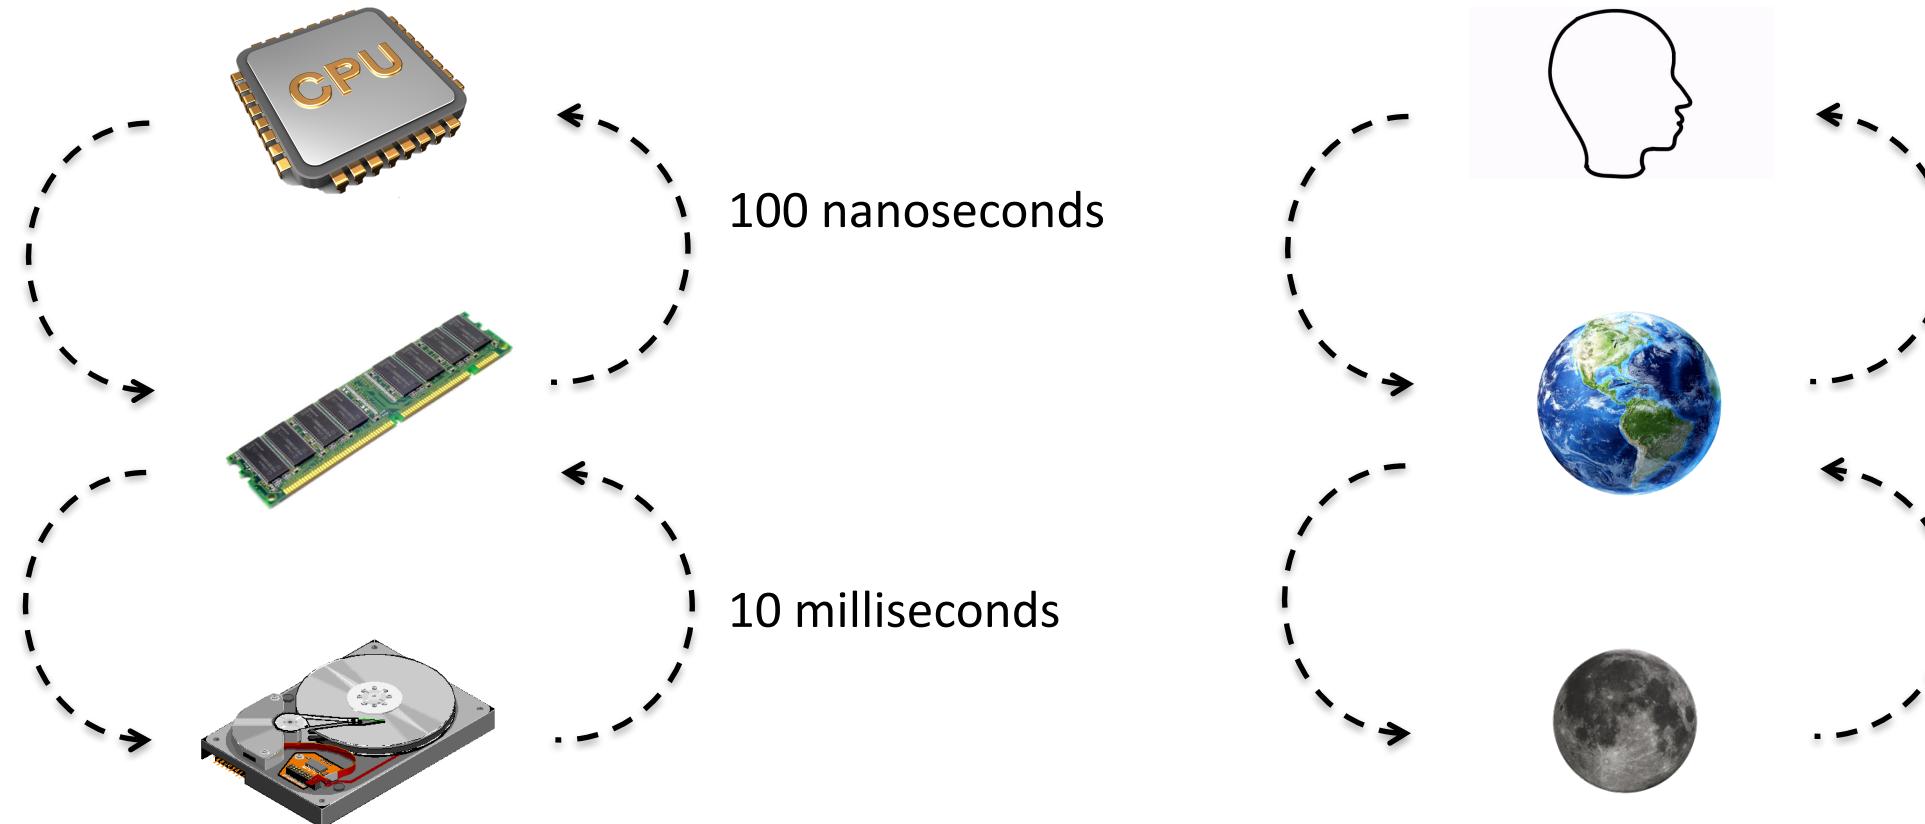
Disk is a bottleneck



Disk is a bottleneck

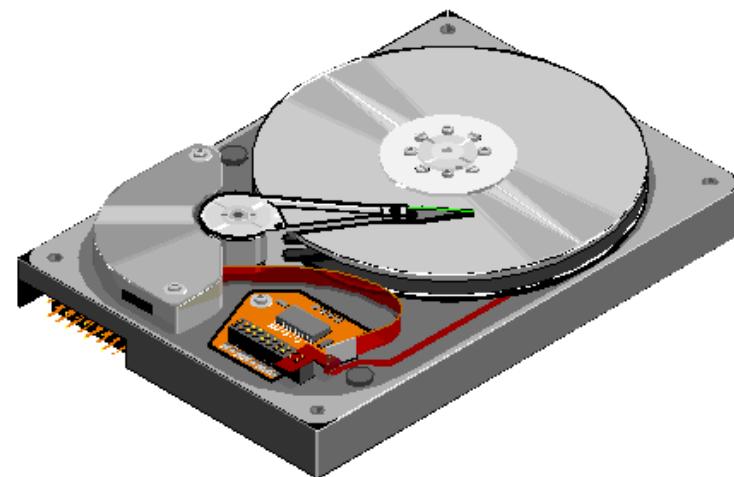


Disk is a bottleneck

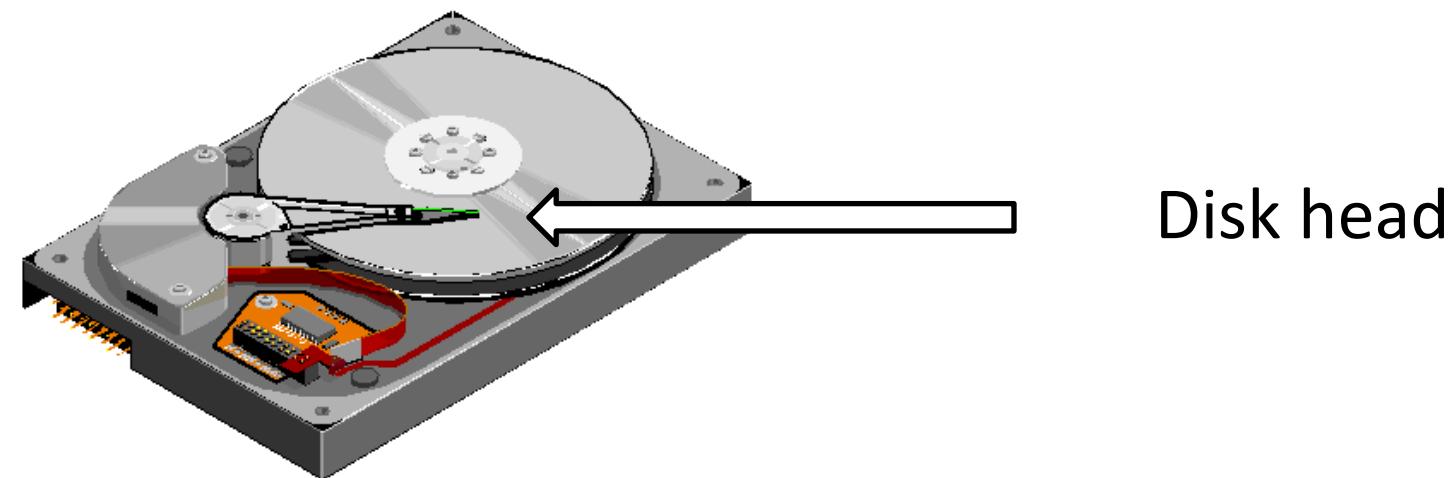


Goal: alleviate disk as the bottleneck

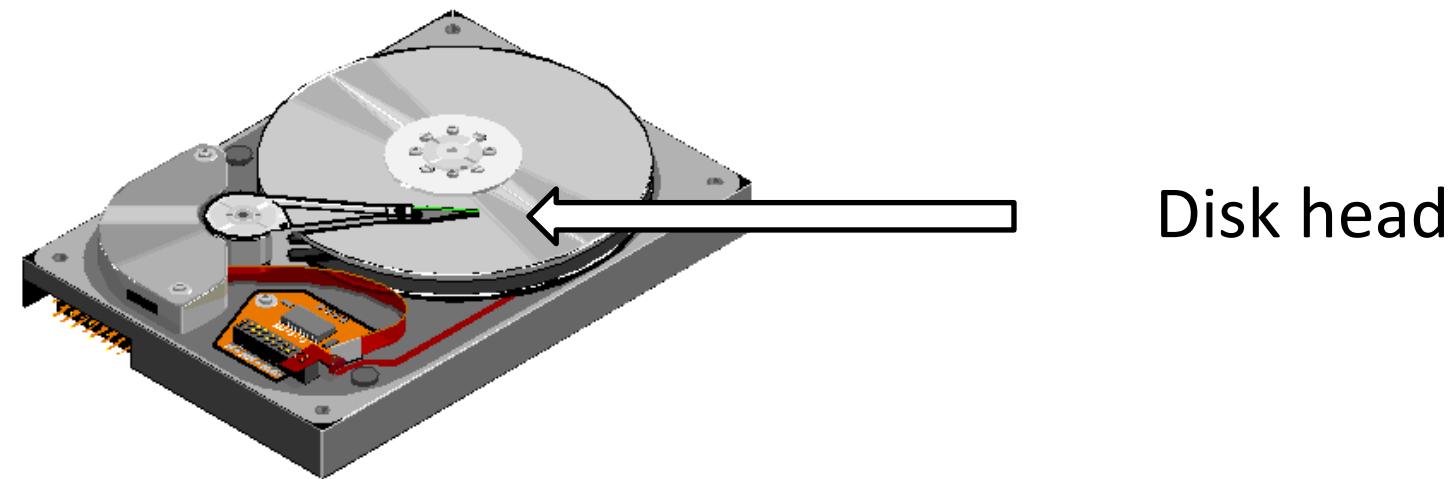
Why is disk slow?



Why is disk slow?



Why is disk slow?



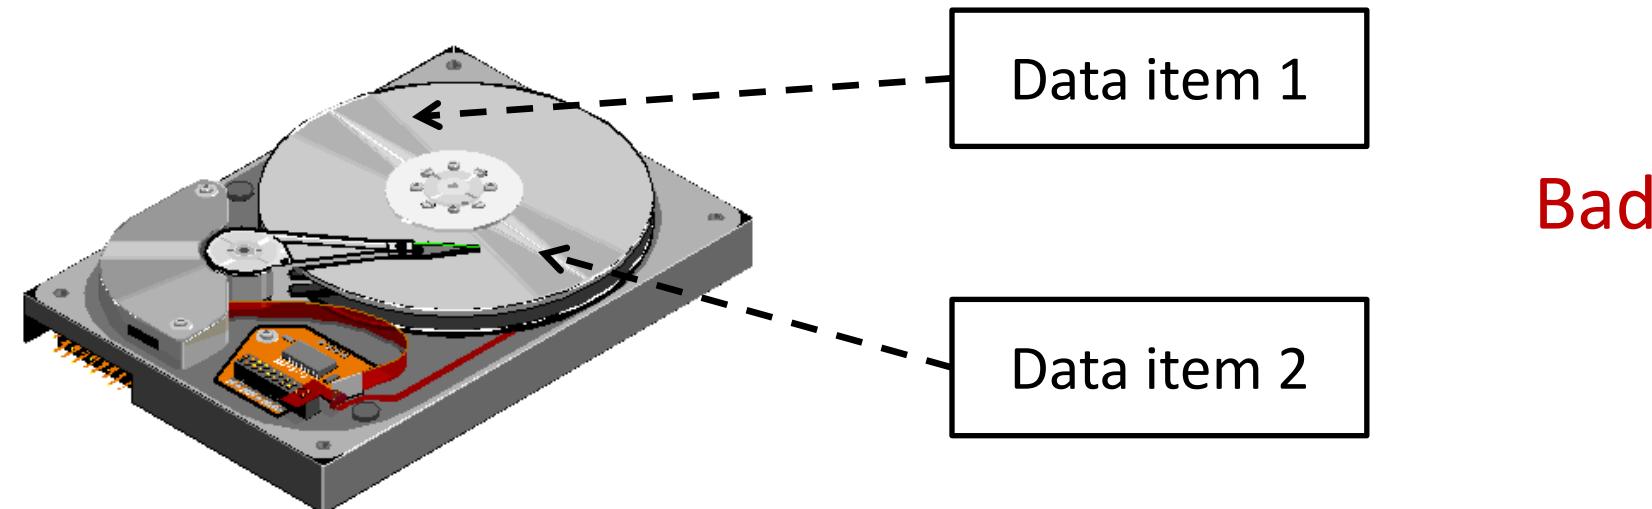
Disk head

Question 1: How to minimize disk access?

Question 2: What to do during a disk access?

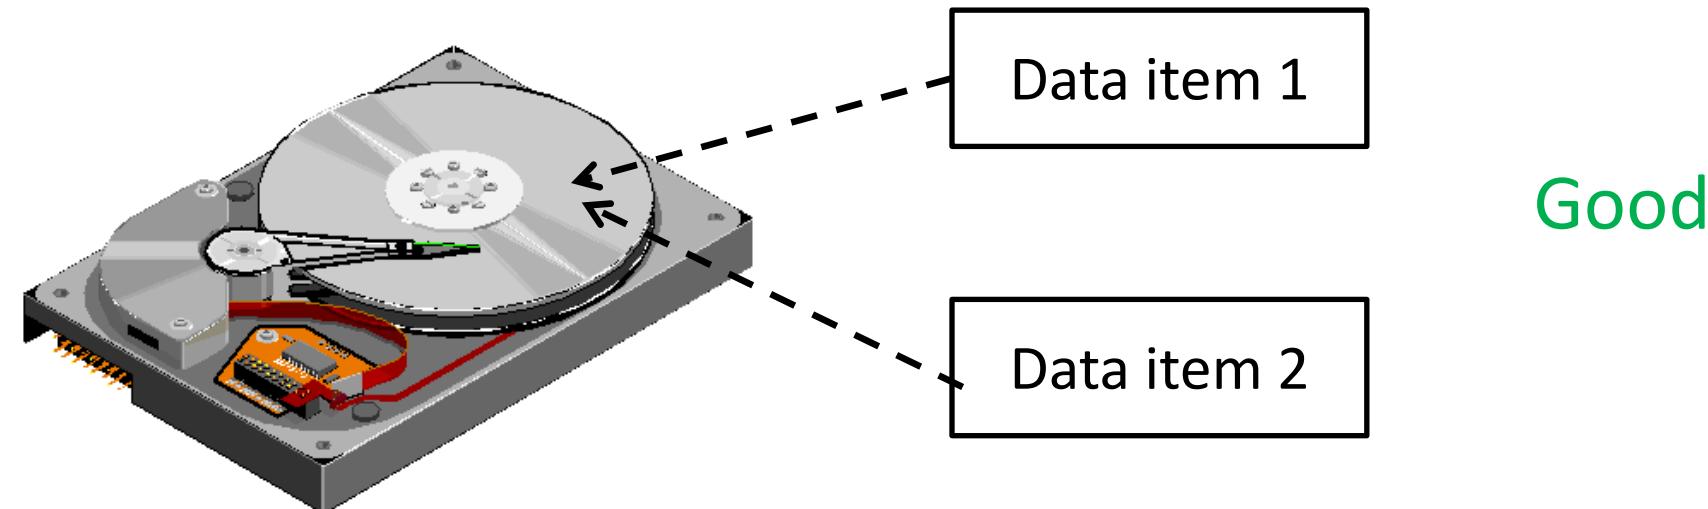
Minimizing Disk Access

Store together data that is frequently co-accessed



Minimizing Disk Access

Store together data that is frequently co-accessed



Row-Store

Co-locate all information about data objects in rows

Example: Bank customer Sara deposits \$100

Main
Memory



Disk



2 disk accesses, since data about
sara is co-located

2

Sara

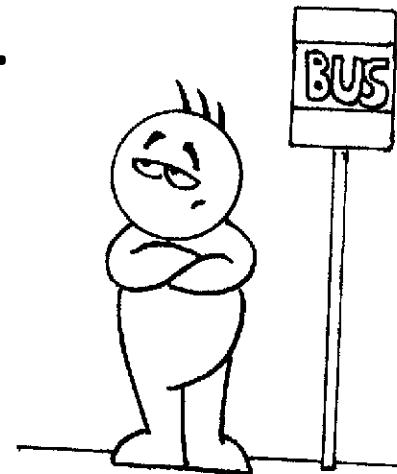
200

Add
100

ID	Name	Balance
1	Bob	100
3	Trudy	450

What to do during a disk access?

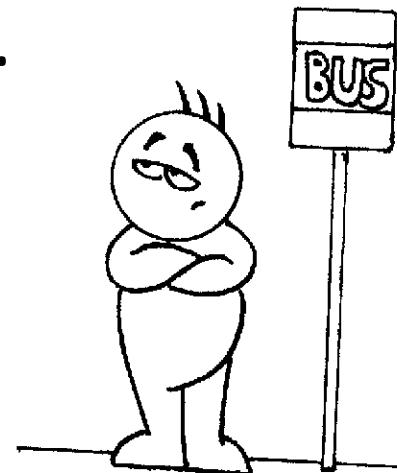
Don't wait.



Start running the next operation(s)

What to do during a disk access?

Don't wait.



Start running the next operation(s)

Improves performance 

But data can get corrupted 

Example of Data Corruption

Bob and Sara share a bank account

Both deposit \$100 at same time

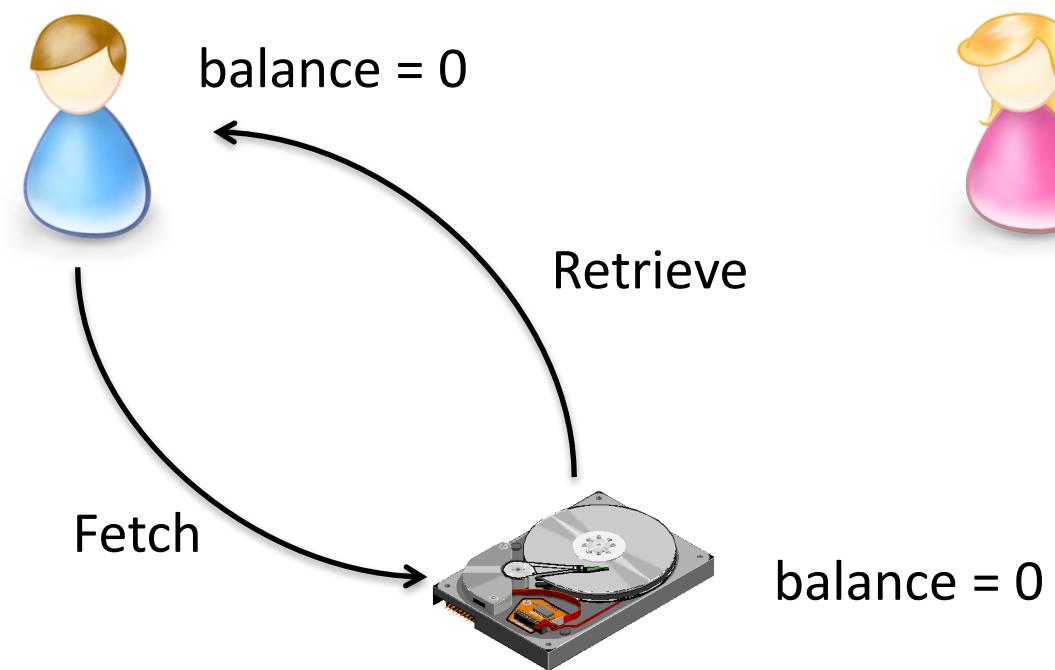


balance = 0

Example of Data Corruption

Bob and Sara share a bank account

Both deposit \$100 at same time



Example of Data Corruption

Bob and Sara share a bank account

Both deposit \$100 at same time



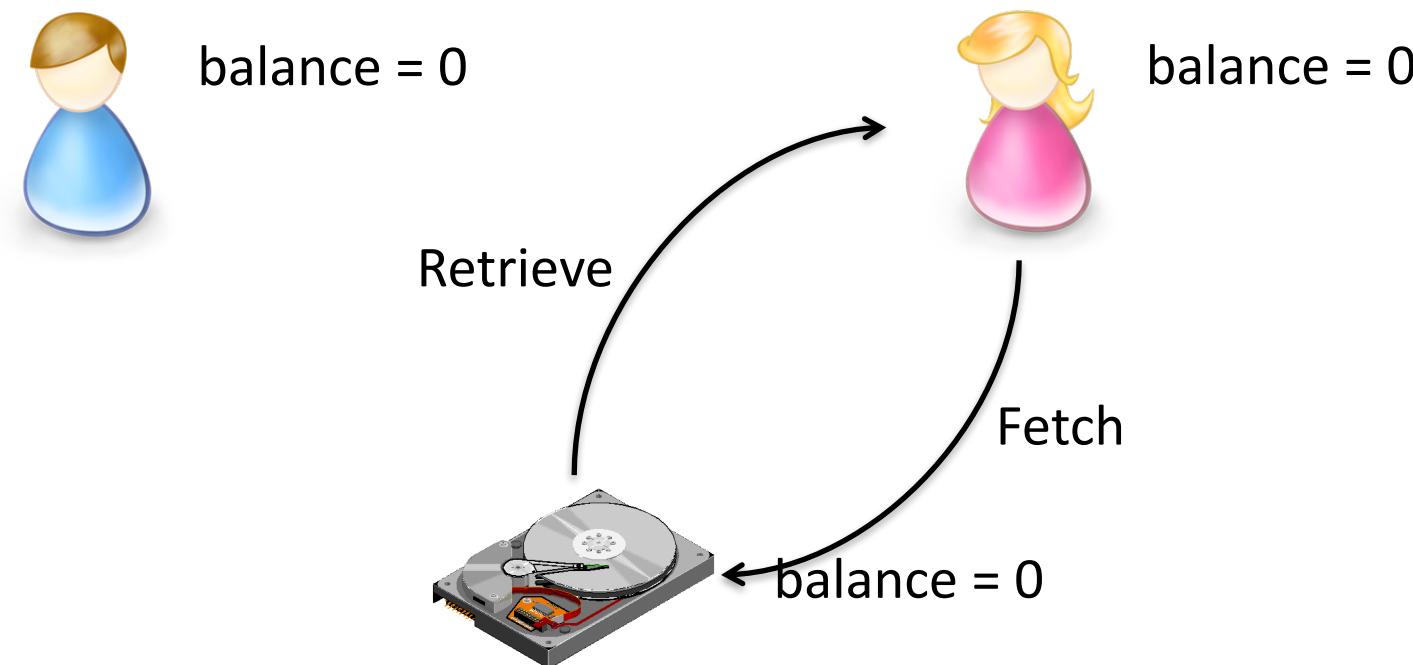
balance = 0



balance = 0

Example of Data Corruption

Bob and Sara share a bank account
Both deposit \$100 at same time



Example of Data Corruption

Bob and Sara share a bank account

Both deposit \$100 at same time



balance = 0



balance = 0



balance = 0

Example of Data Corruption

Bob and Sara share a bank account

Both deposit \$100 at same time



balance = 100



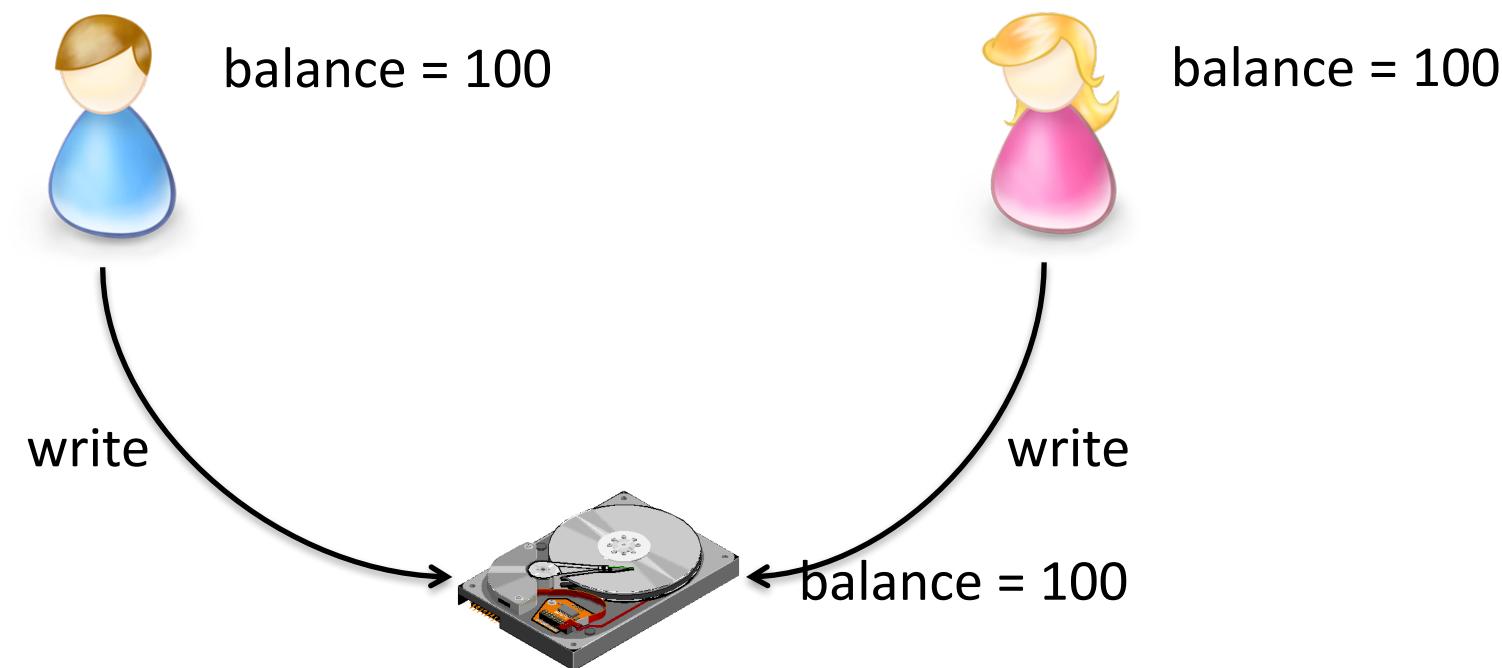
balance = 100



balance = 0

Example of Data Corruption

Bob and Sara share a bank account
Both deposit \$100 at same time



Example of Data Corruption

Bob and Sara share a bank account

Both deposit \$100 at same time



balance = 100

Example of Data Corruption

Bob and Sara share a bank account

Both deposit \$100 at same time



Account balance should be 200!
Bob and Sara lost money.

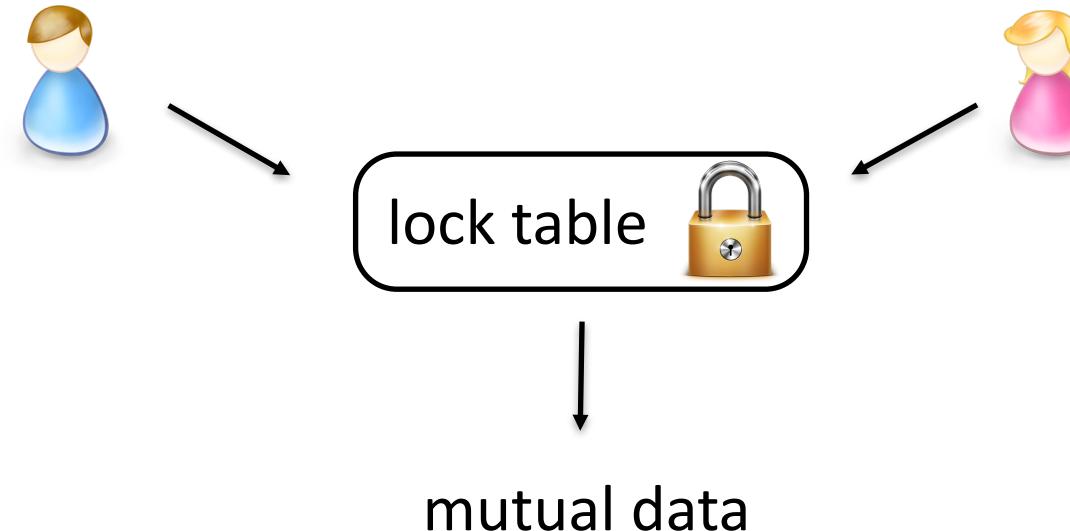


balance = 100

How to avoid data corruption?

Insight: operations can be concurrent, as long as they don't modify the same data

Solution: locking



Database Design Goals

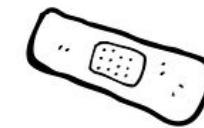
1. **Speed**



2. Affordability



3. Resilience to system failure



Database Design Goals

1. Speed



2. **Affordability**



3. Resilience to system failure



Disk Was Expensive

1 gigabyte for \$10000 in 1980

Avoid storing replicas of same data



ID	name	account-ID	balance
1	Bob	1	100
2	Sara	1	100
3	Trudy	2	450

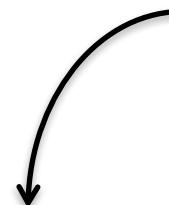
Disk Was Expensive

1 gigabyte for \$10000 in 1980

Avoid storing replicas of same data

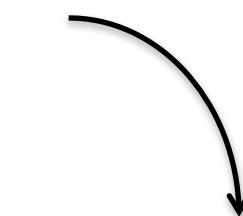


ID	name	account-ID	balance
1	Bob	1	100
2	Sara	1	100
3	Trudy	2	450



Customers

ID	name	account-ID
1	Bob	1
2	Sara	1
3	Trudy	2



Accounts

ID	balance
1	100
2	450

Normalization

Saves storage space



Helps data integrity



Reads take longer



Customers

ID	name	account-ID
1	Bob	1
2	Sara	1
3	Trudy	2

Accounts

ID	balance
1	100
2	450

SQL

Getting Bob's account balance:

```
select balance  
from Customers c, Accounts a  
where c.account-ID = a.ID and c.name = "Bob"
```

Customers

ID	name	account-ID
1	Bob	1
2	Sara	1
3	Trudy	2

Accounts

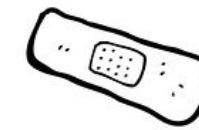
ID	balance
1	100
2	450

Database Design Goals

1. Speed
2. **Affordability**
3. Resilience to system failure



\$



Database Design Goals

1. Speed



2. Affordability

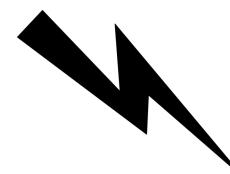
\$

3. **Resilience to system failure**

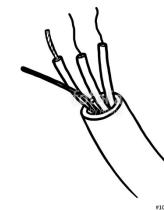


Possible Failures

Power failure



Hardware failure



Natural disaster



Data is precious (e.g. bank)

Provide recovery mechanism

Power Failure

Sara transfers \$100 to Trudy



Sara's
balance = 0



Sara's
balance = 100

Trudy's
balance = 450



Power Failure

Sara transfers \$100 to Trudy

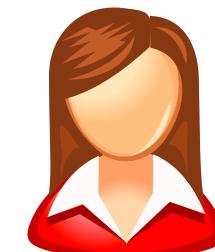


Sara's
balance = 0



Sara's
balance = 100

Trudy's
balance = 450



Power Failure

Sara transfers \$100 to Trudy



Sara's
balance = 0

Trudy's
balance = 550

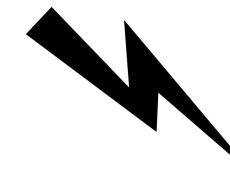
Trudy's
balance = 450



Power Failure

Sara transfers \$100 to Trudy

power fails



Sara's
balance = 0

Trudy's
balance = 550



Trudy's
balance = 450

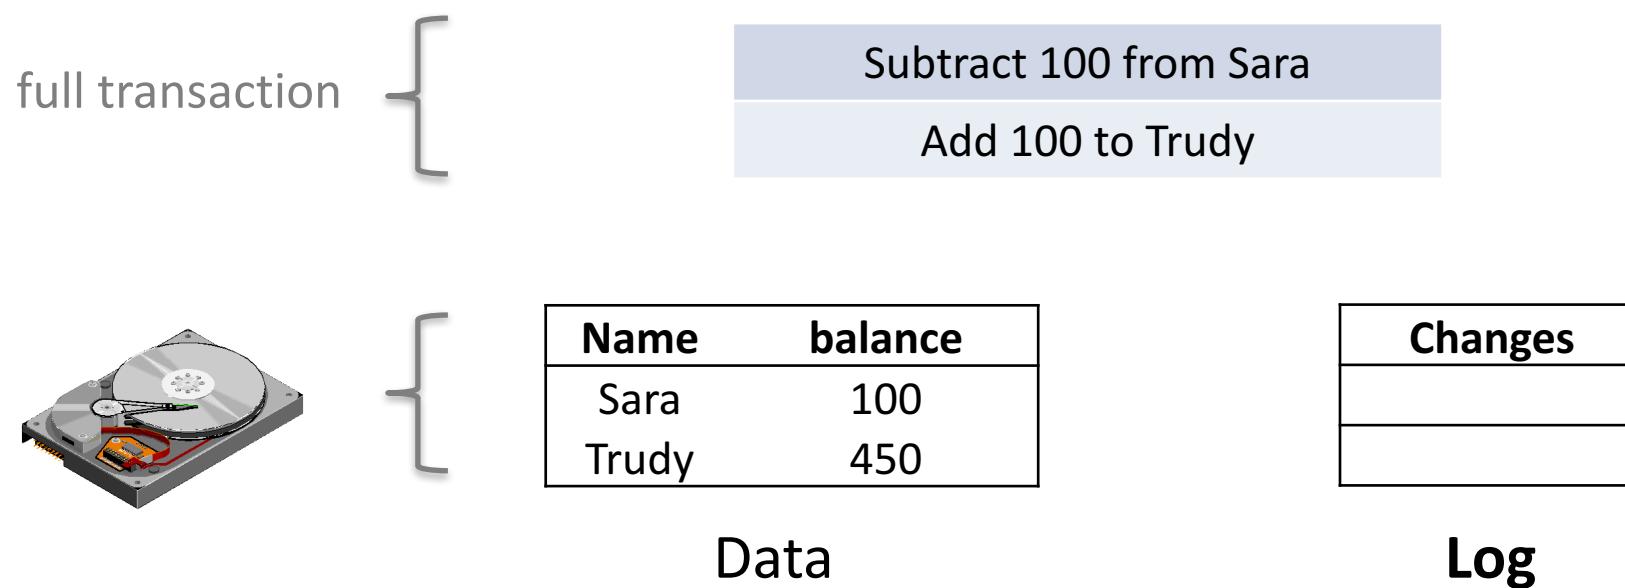
Transactions

Atomic sequence of operations (all or nothing)



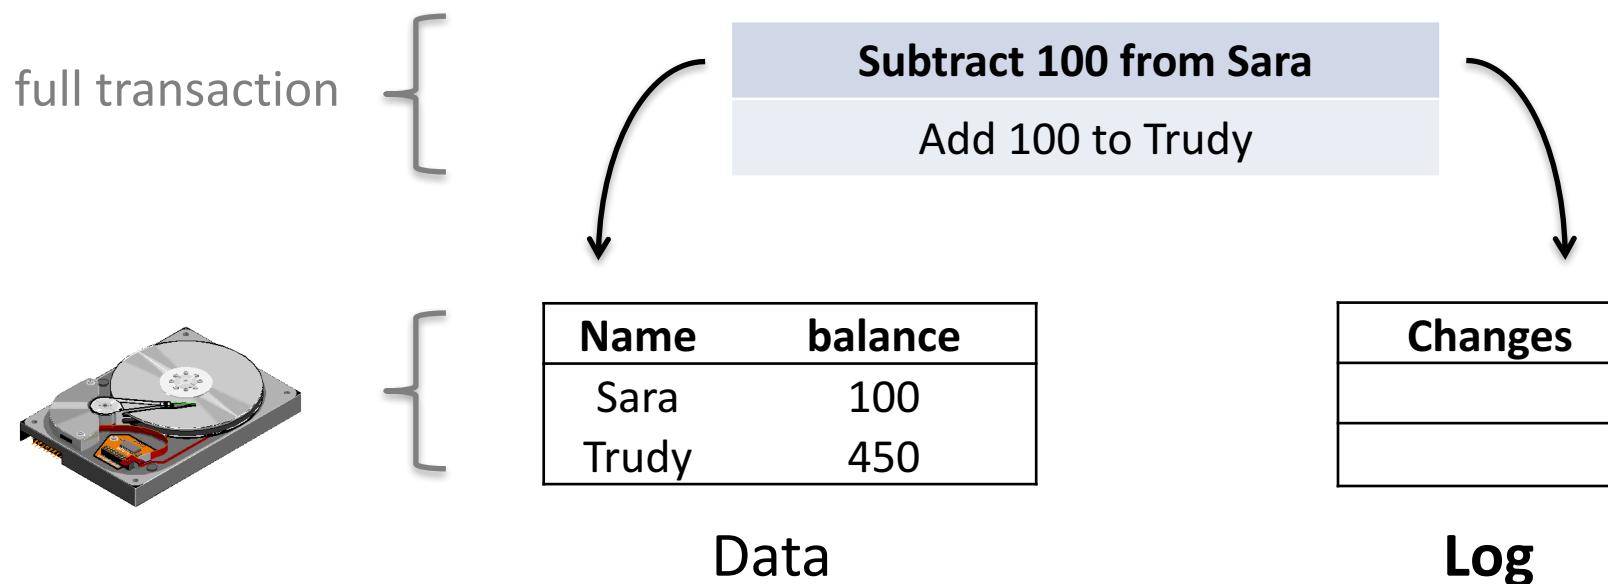
Transactions

Atomic sequence of operations (all or nothing)



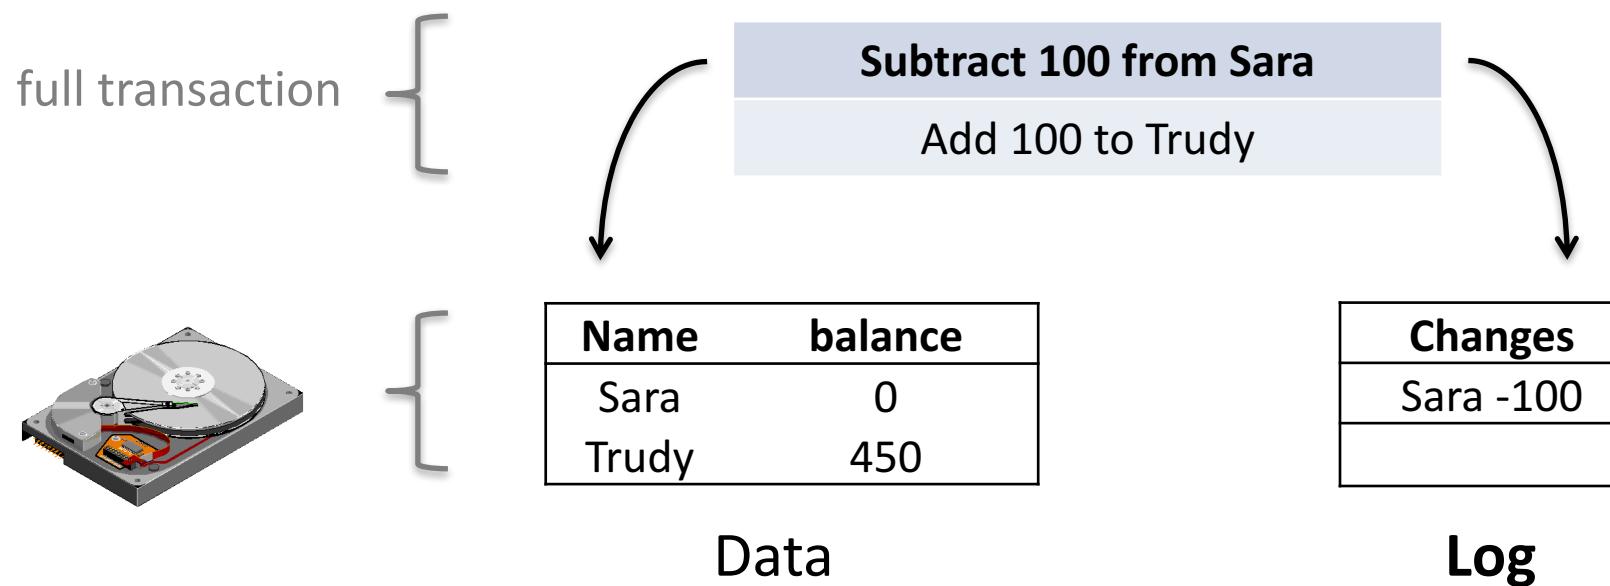
Transactions

Atomic sequence of operations (all or nothing)



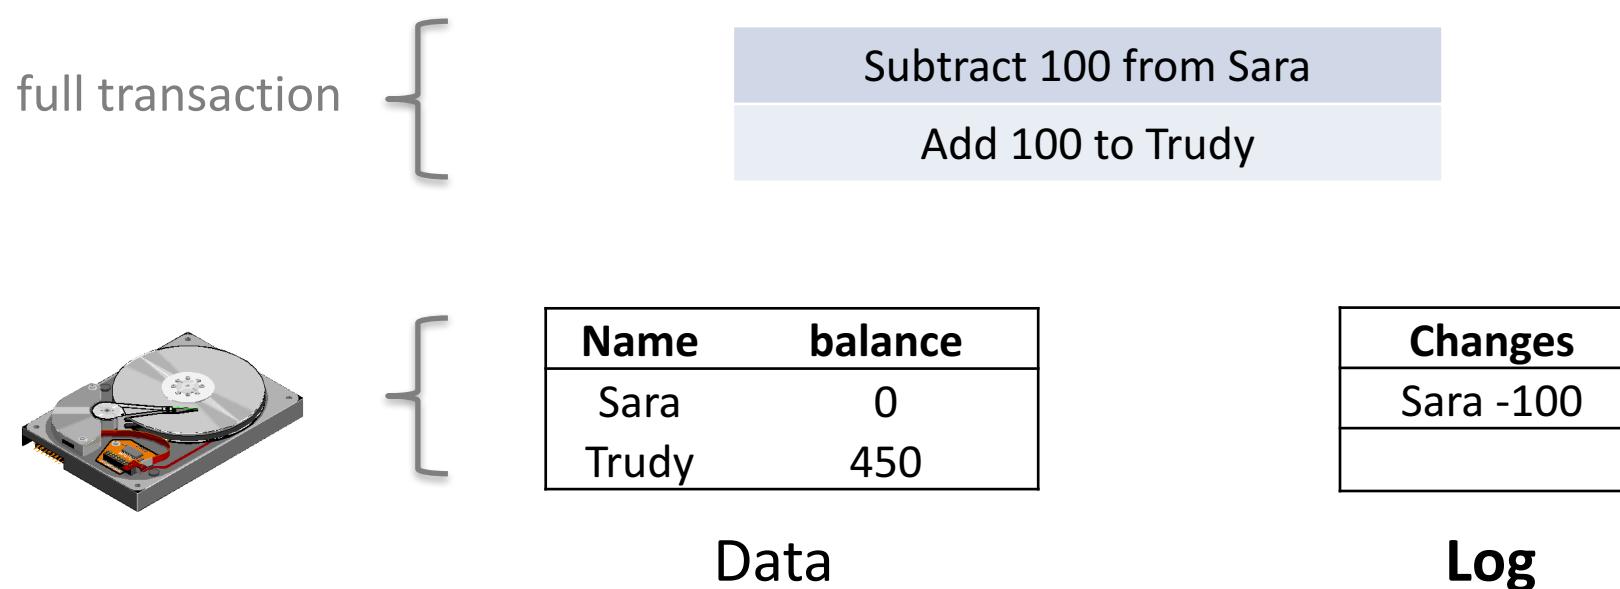
Transactions

Atomic sequence of operations (all or nothing)



Transactions

Atomic sequence of operations (all or nothing)



Transactions

Atomic sequence of operations (all or nothing)

full transaction {

Subtract 100 from Sara
Add 100 to Trudy



{

Name	balance
Sara	0
Trudy	450

Data

Changes
Sara -100

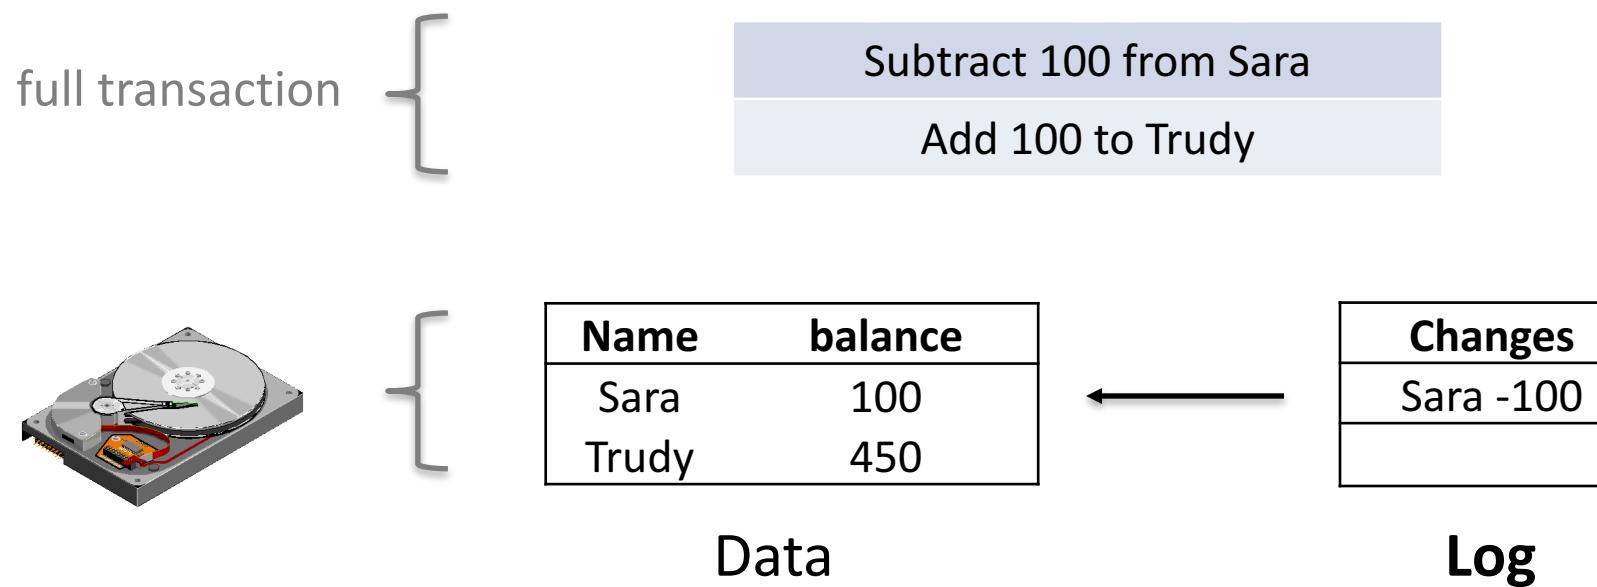
Log

After failure, undo changes by unfinished transactions



Transactions

Atomic sequence of operations (all or nothing)



ACID



Data integrity pitfalls

Concurrency (fix with locking)

Power failure (fix with logging)

Transactions should be

Atomic

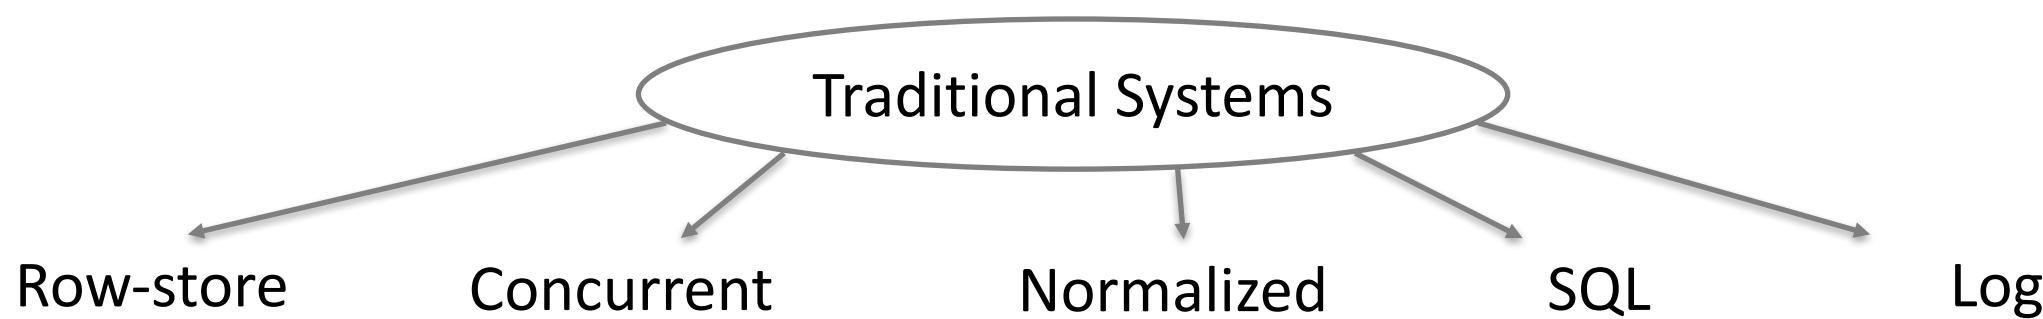
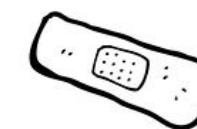
Consistent

Isolated

Durable

Database Design Goals

1. Speed
2. Affordability
3. Resilience to system failure

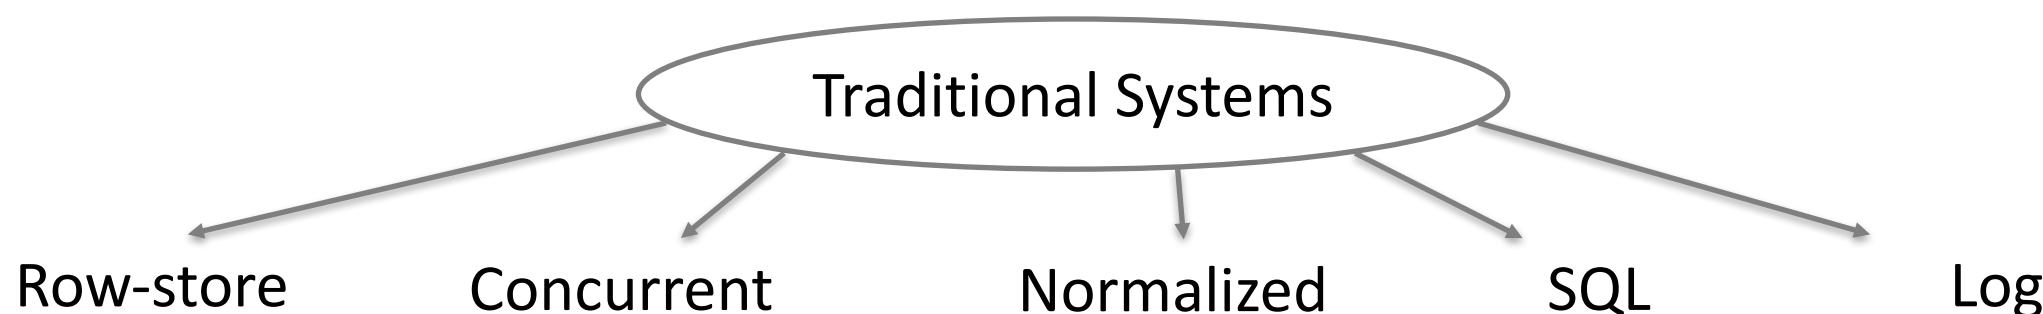
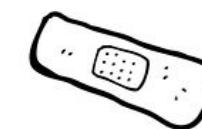


Database Design Goals

1. Speed
2. Affordability
3. Resilience to system failure



\$



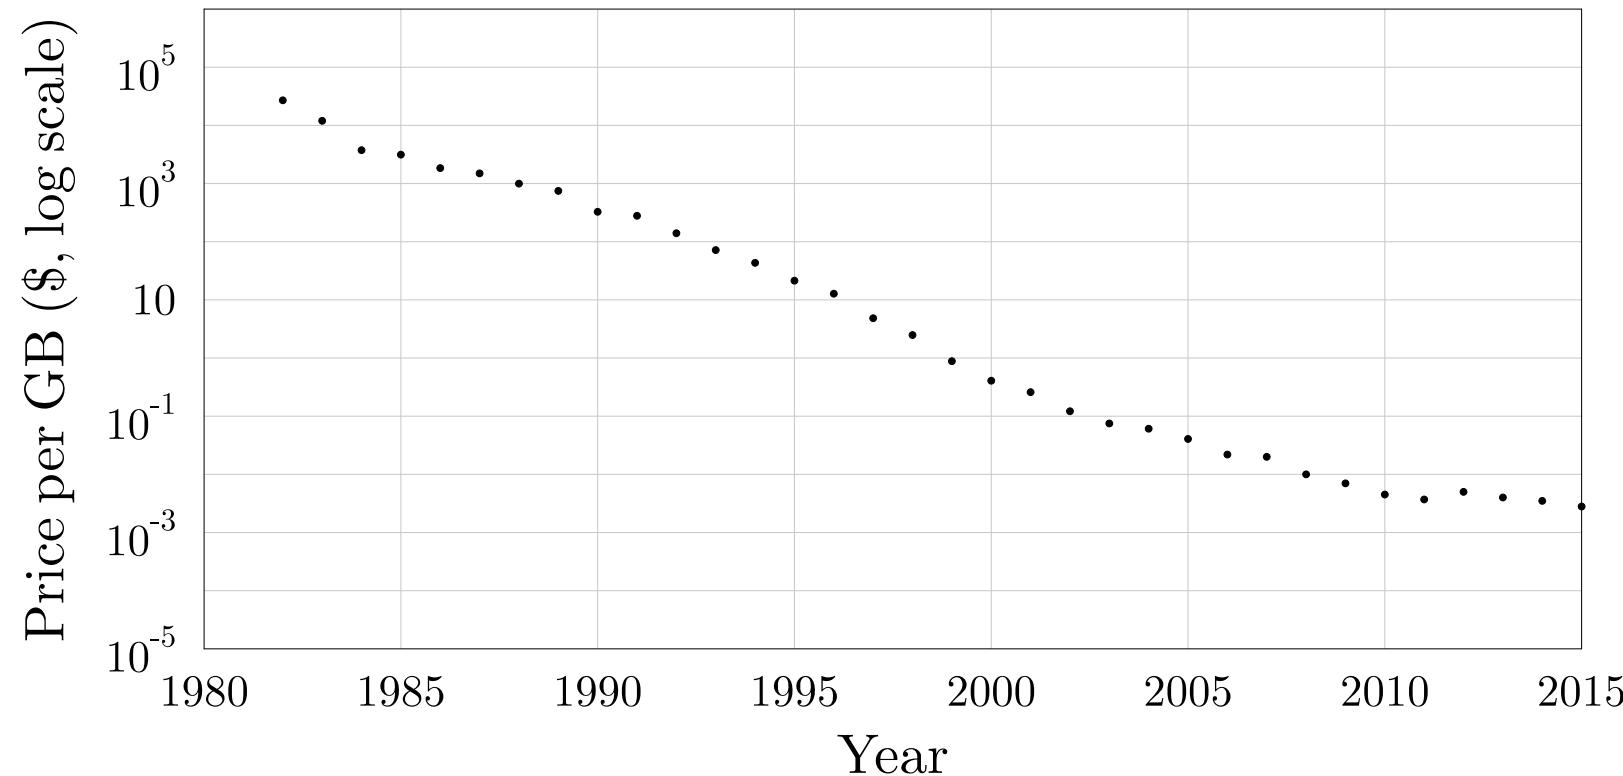
Everything is Motivated by hardware.

BigData Challenges

What changed in hardware?

Why can't traditional systems cope?

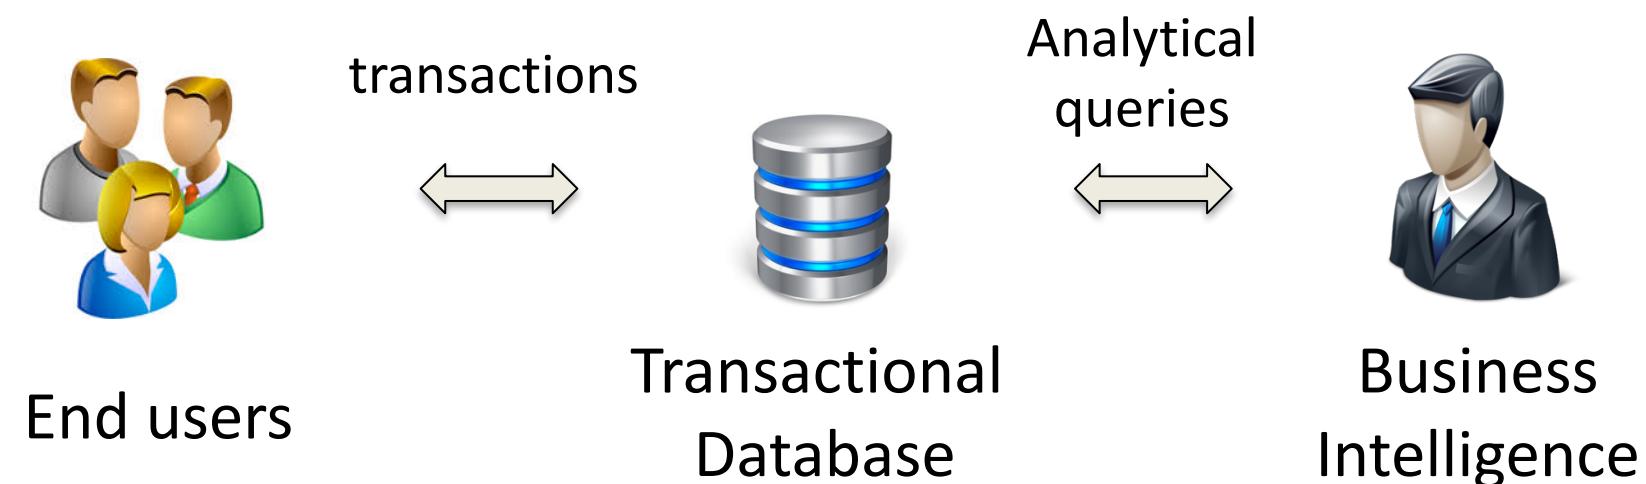
Declining Storage Costs



$10^7 \times$ cheaper

Problem 1: Analyzing Big Data

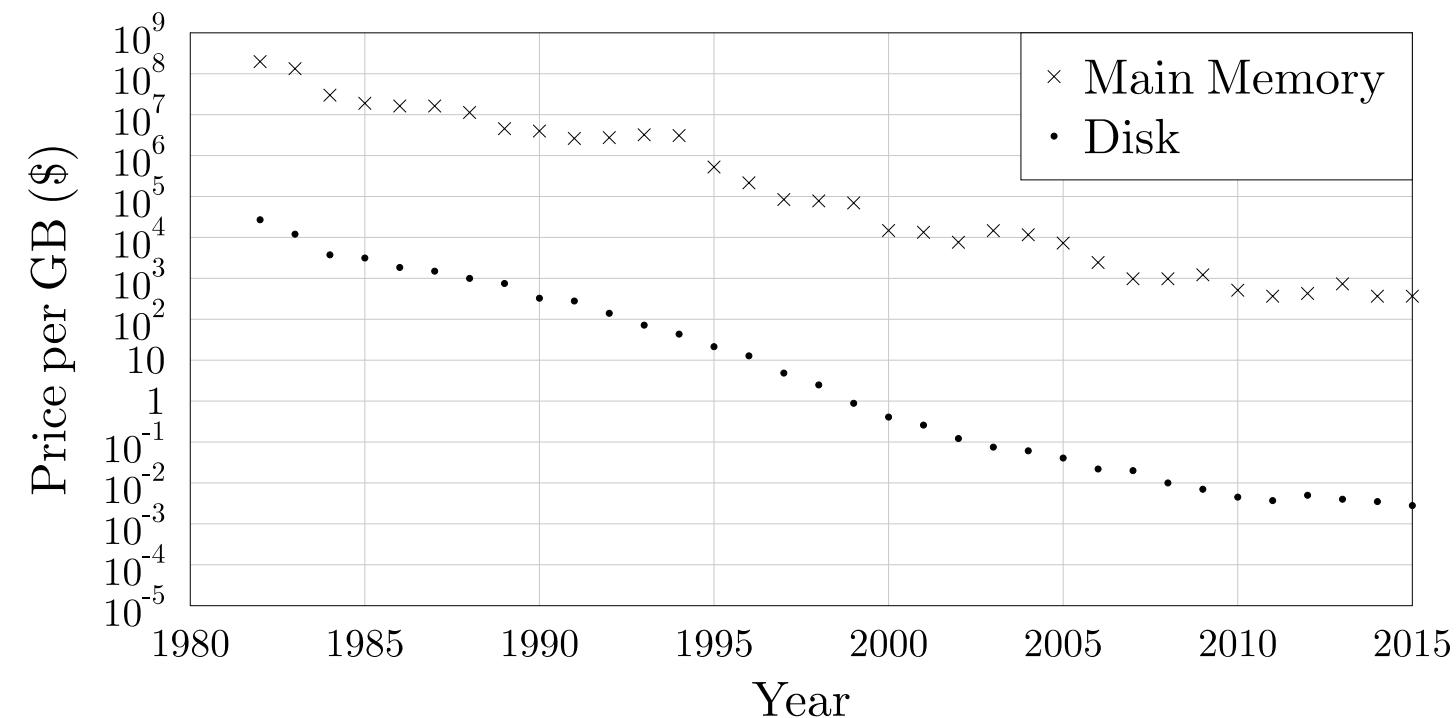
Everyone is hording data



Analytical queries are expensive

Worsens as data grows

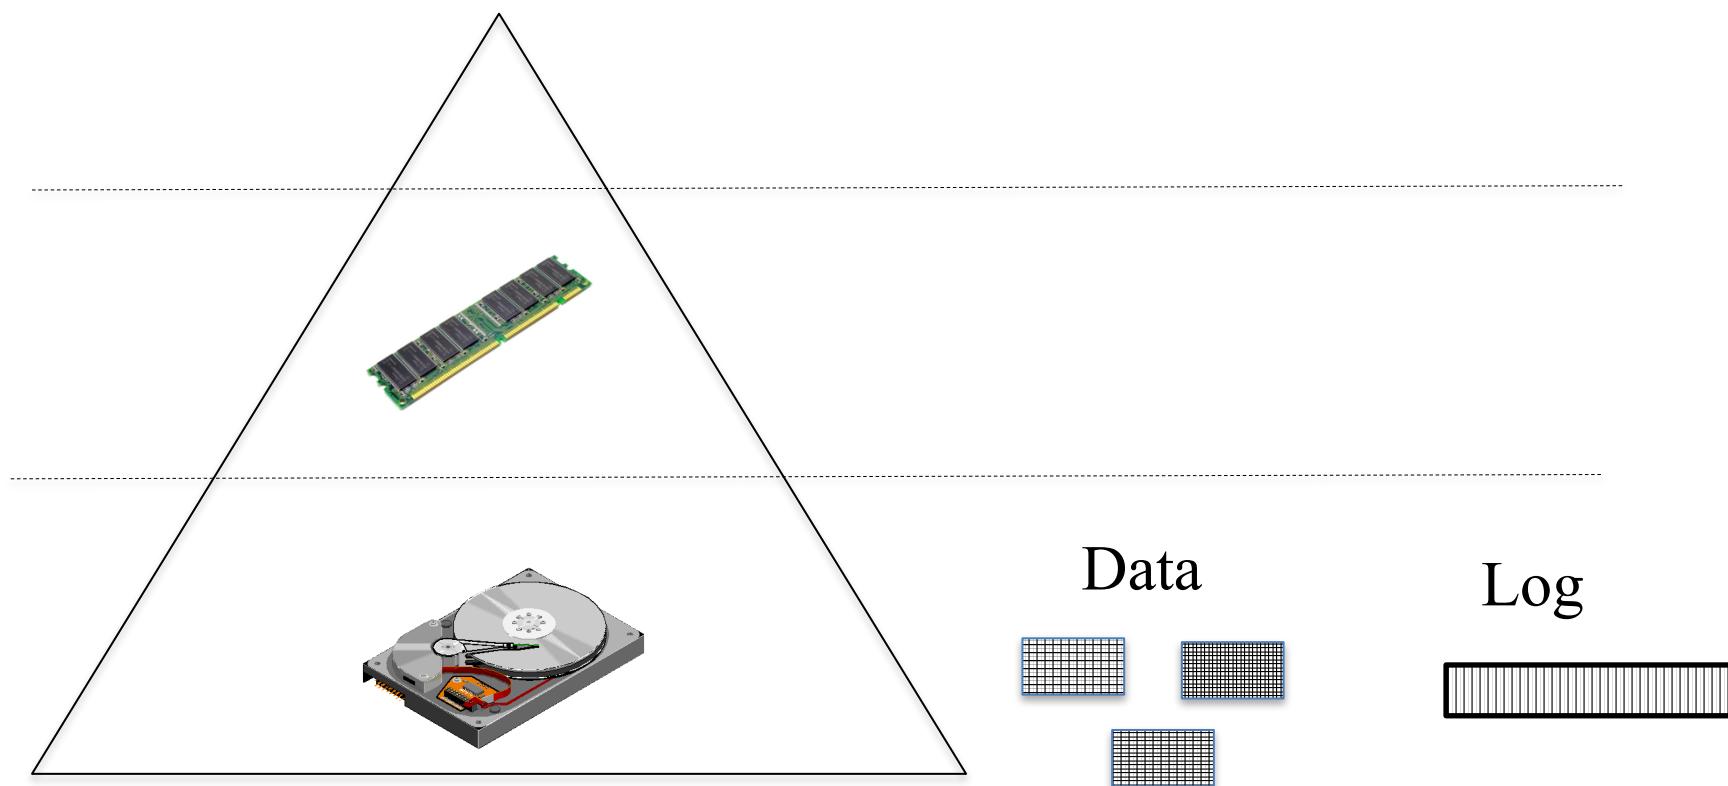
Declining Main Memory Costs



Main memory: $10^6 \times$ cheaper

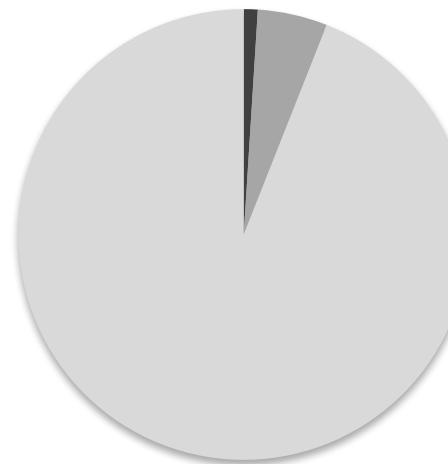
Disk: $10^7 \times$ cheaper

Some Databases Fit in Memory



Problem 2: ACID is a new Bottleneck

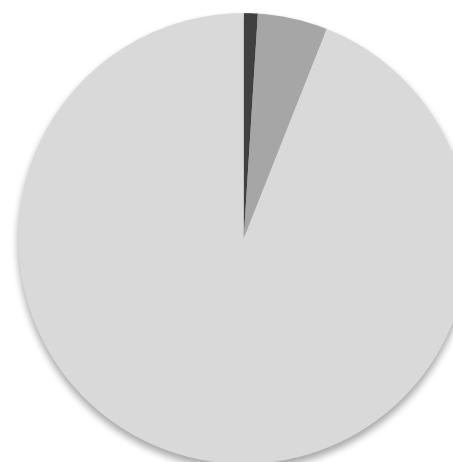
Disk-based system



- Useful work
- ACID
- Disk access

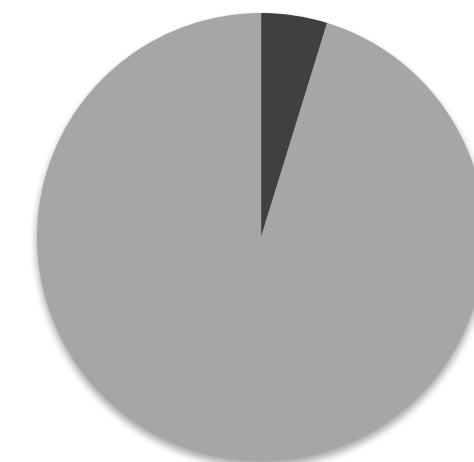
Problem 2: ACID is a new Bottleneck

Disk-based system



- Useful work
- ACID
- Disk access

Memory-based system



- Useful work
- ACID

Locking, logging

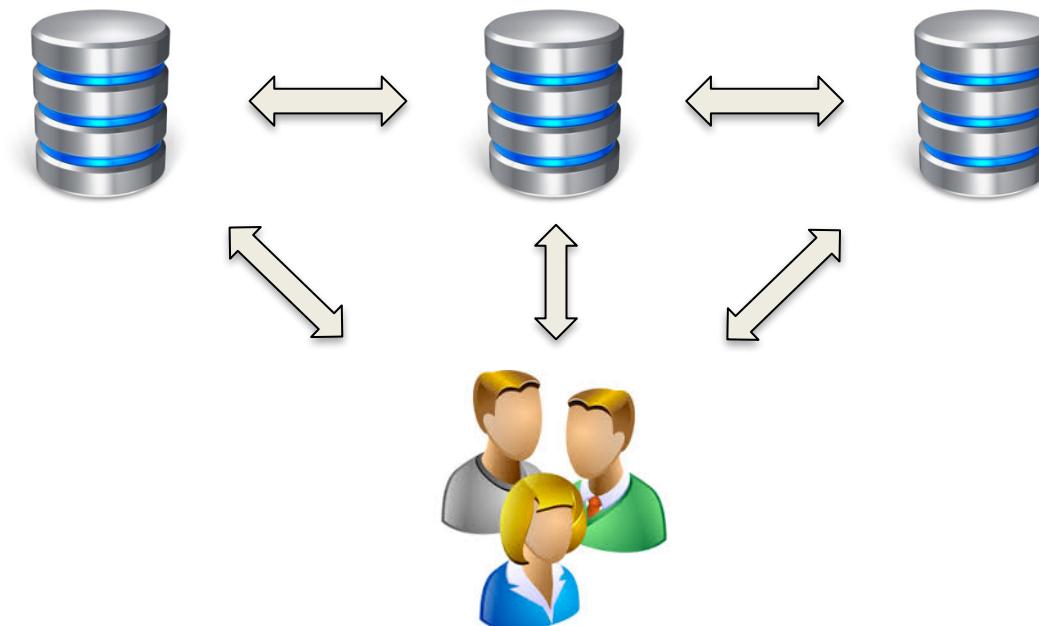
Distribution Across Machines

Increase bandwidth

Remain available if one machine fails

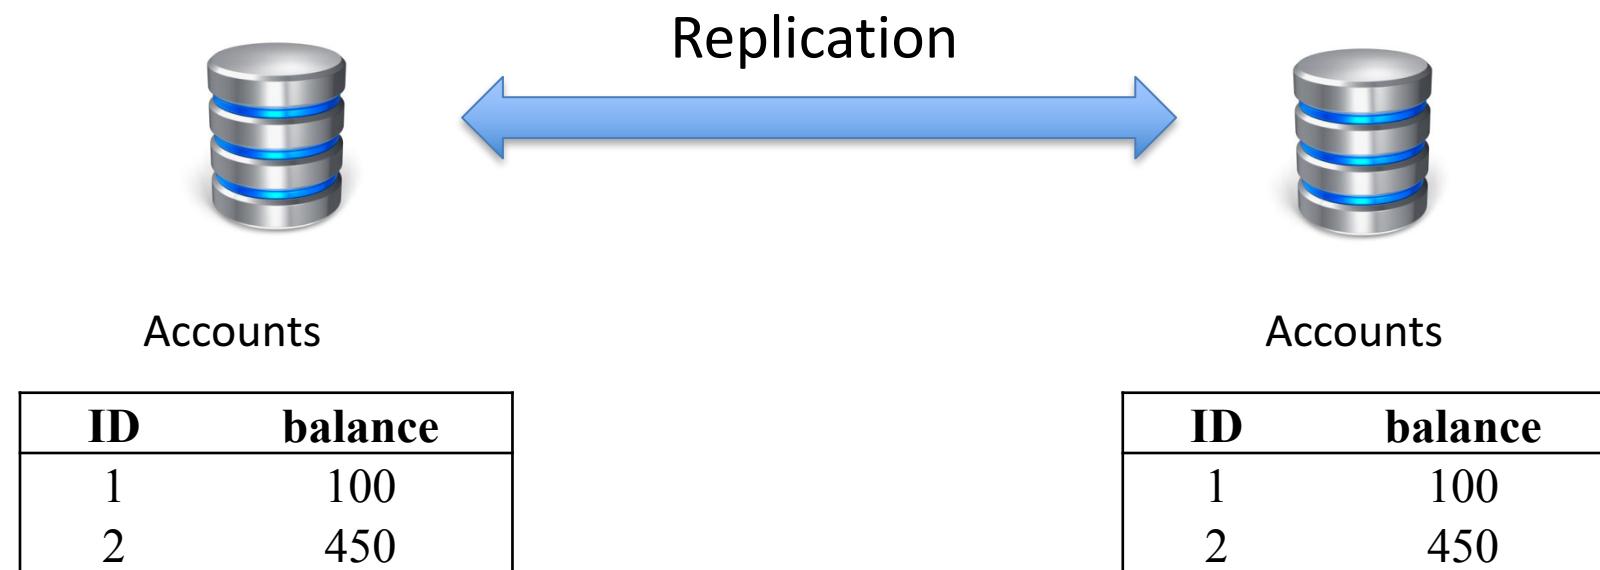
(partitioning)

(replicating)



Problem 3: Consistency

Updates require synching across network



Problem 4: Joins

Take longer as tables grow

E.g. find Bob's account balance

Customers

ID	name	account-ID
	...	
2	Bob	X
	...	

Accounts

ID	balance
	...
X	100
	...

Problems

1. Analytical Queries
2. ACID
3. Consistency
4. Joins

Problems

1. Analytical Queries
2. ACID
3. Consistency
4. Joins

Solutions

- Analytical Databases
- NoSQL Databases
- NewSQL Databases

Problems

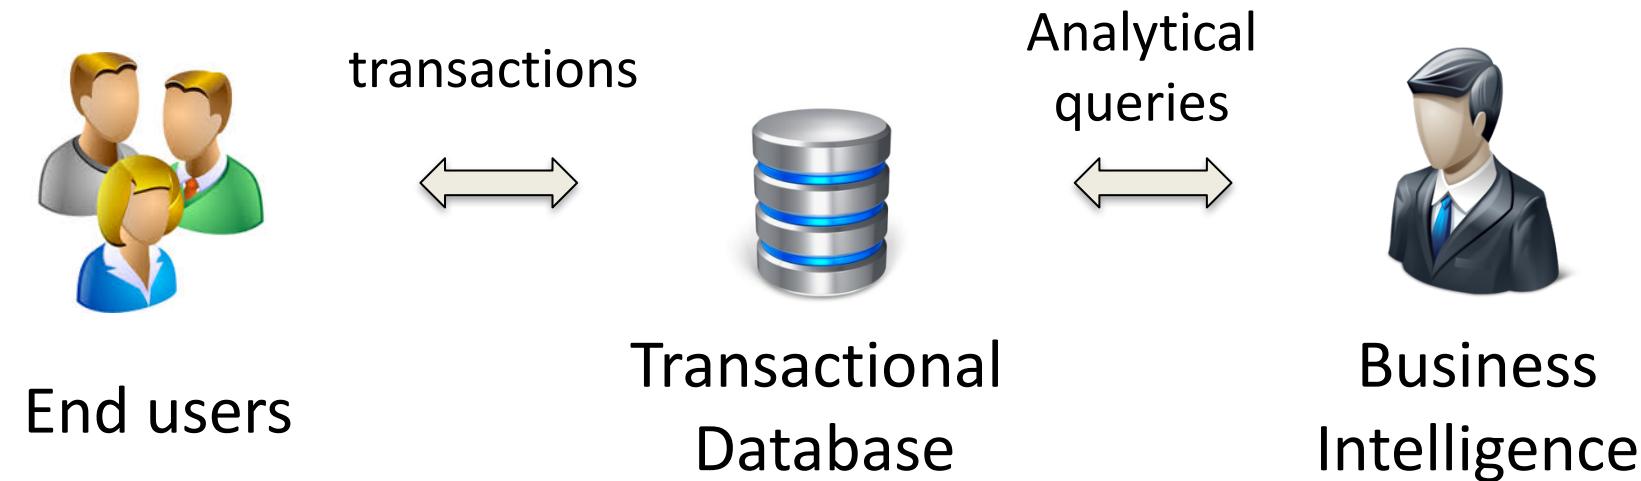
1. Analytical Queries
2. ACID
3. Consistency
4. Joins

Solutions

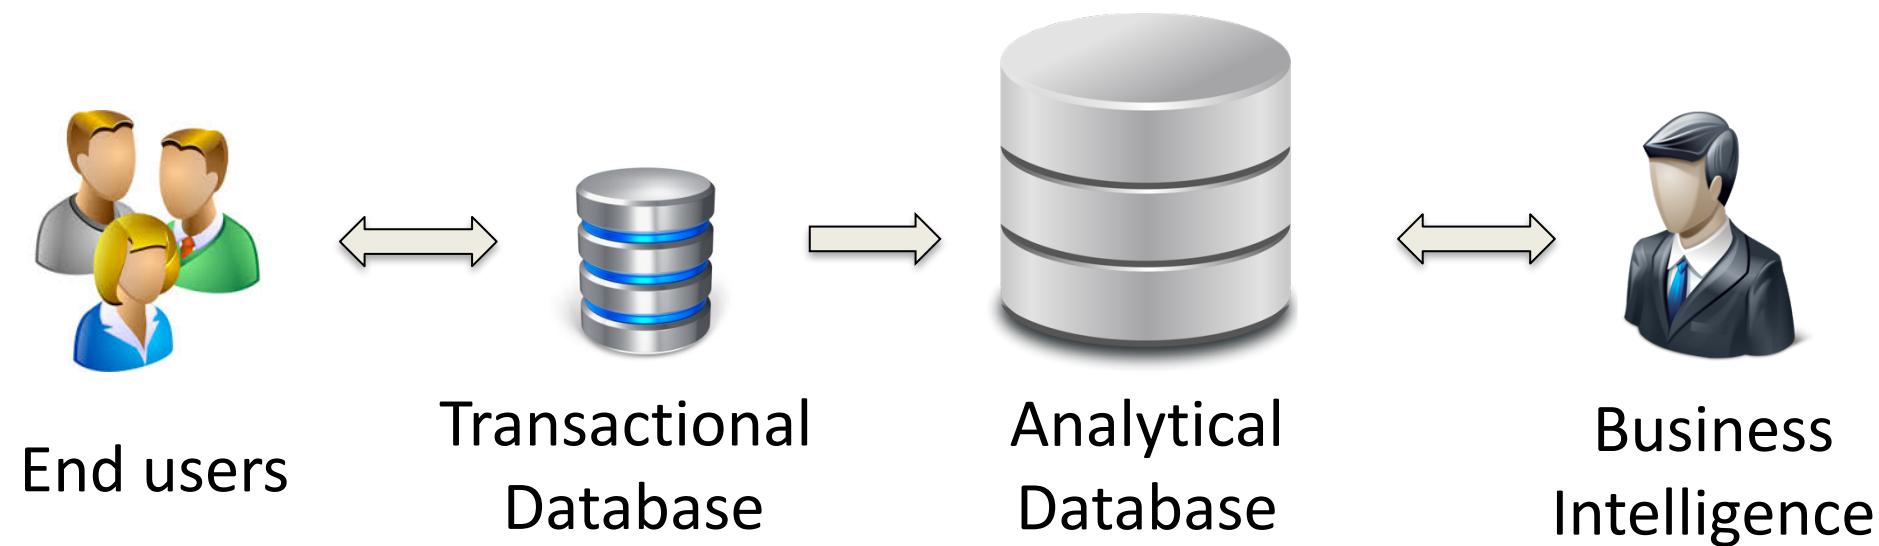


Analytical Databases
NoSQL Databases
NewSQL Databases

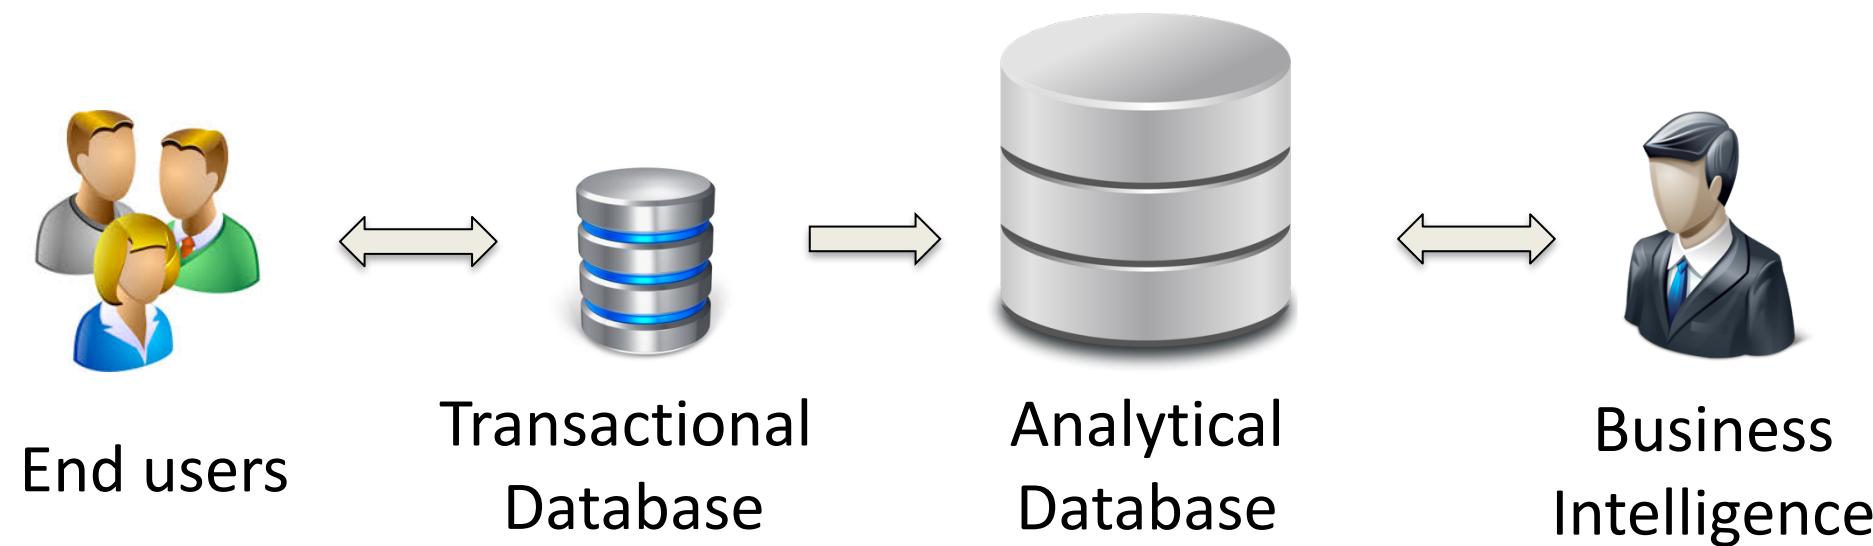
Analytical Databases



Analytical Databases



Analytical Databases



Different internal design for analytics

Analytical Queries Examples

order id	cust id	product id	price	order date	receipt date	priority	status	comment
...

How long is delivery?

Revenue from product X?

Few columns, many rows

Problem: Data is stored row by row

Column-Stores

order id	cust id	status	price	order date	receipt date	priority	clerk	comment
...

Each column is stored separately

Changes entire architecture

Examples: Vertica, Vectorwise, etc.



Problems

1. Analytical Queries
2. ACID
3. Consistency
4. Joins

Solutions

- 
- Analytical Databases
 - NoSQL Databases
 - NewSQL Databases

Problems

1. Analytical Queries
2. ACID
3. Consistency
4. Joins

Solutions

- Analytical Databases
- NoSQL Databases
- NewSQL Databases

NoSQL

Named in 2009

Hashtag to advertise a conference

NoSQL

Problems

1. Joins
2. ACID
3. Consistency

NoSQL

Problems

1. Joins
2. ACID
3. Consistency



De-normalization

Normalization Recap

Reduce data volume



Introduces joins



Locking & logging across

People

ID	name	addr-ID
1	Sara	1
2	Bob	1
3	Will	2

Addresses

ID	addr
1	33 Oxford st
2	45 Harvard st

Normalization Recap

-  Reduce data volume smaller issue
-  Introduces joins bigger issue
-  Locking & logging across bigger issue

People

ID	name	addr-ID
1	Sara	1
2	Bob	1
3	Will	2

Addresses

ID	addr
1	33 Oxford st
2	45 Harvard st

Normalization Recap

-  Reduce data volume smaller issue
-  Introduces joins bigger issue
-  Locking & logging across bigger issue

ID	name	addr-ID
1	Sara	1
2	Bob	1
3	Will	2

Addresses	
ID	addr
1	33 Oxford st
2	45 Harvard st

NoSQL
→

ID	name	Addr-ID
1	Sara	33 Oxford st
2	Bob	33 Oxford st
3	Will	45 Harvard st

Document-store

Each object is a document. Flexible format.

```
{ name: "Bob",
  balance: 100,
  favorite-color: "red"
  credit-score: 3.0
}
```

```
{ name: "Sara",
  balance: 100
  hobbies: ["rowing", "running"]
}
```

```
db.customers.find(name:"Sara")
```

Downside: costly to search a document

Transaction spans Single Document

Cannot maintain integrity guarantees for changes across documents

Non-interacting objects

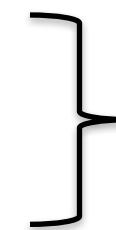
No sharing (e.g. bank account)

No exchanging (e.g. money transfers)

NoSQL

Problems

1. Joins
2. ACID
3. Consistency



De-normalization

NoSQL

Problems

1. Joins
 2. ACID
 3. **Consistency**
- 
- De-normalization
- Relax

Consistency

Different operation order across replicas



Deposit
Add 100

```
{ ...  
    balance: 100  
... }
```



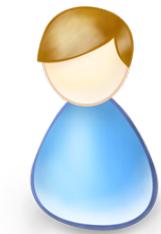
Interest
Multiply by 1.1

```
{ ...  
    balance: 100  
... }
```

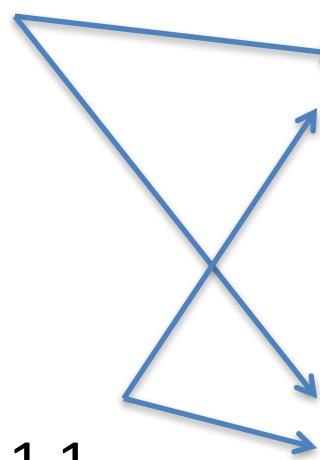


Consistency

Different operation order across replicas



Deposit
Add 100



{ ...
balance: 100
... }



Interest
Multiply by 1.1

{ ...
balance: 100
... }



Consistency

Different operation order across replicas



Deposit
Add 100

Inconsistent
replicas



Interest
Multiply by 1.1

{ ...
balance: 220
... }



{ ...
balance: 210
... }



Consistency

Different operation order across replicas



Deposit
Add 100

Inconsistent
replicas



Interest
Multiply by 1.1

```
{ ...  
    balance: 220  
... }
```



```
{ ...  
    balance: 210  
... }
```

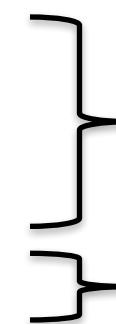


Eventually inconsistencies are fixed.

NoSQL

Problems

1. Joins
2. ACID
3. **Consistency**



De-normalization
Eventual-consistency

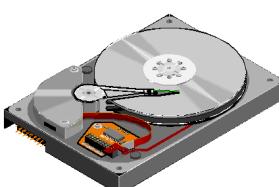
Two NoSQL Types

Main memory



Couchbase

Disk



Two NoSQL Types

Main memory



Couchbase

Disk



My Area!



Niv Dayan

When to use NoSQL?

Non-interacting data objects

Weak consistency

Flexible data model

Quick & dirty & easy

Problems

1. Analytical Queries
2. ACID
3. Consistency
4. Joins

Solutions

- Analytical Databases
- NoSQL Databases
- NewSQL Databases

Problems

1. Analytical Queries
2. ACID
3. Consistency
4. Joins

Solutions

Analytical Databases
NoSQL Databases
NewSQL Databases

Address without compromising generality of application

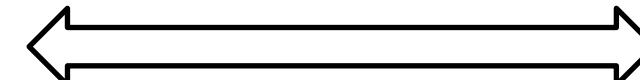
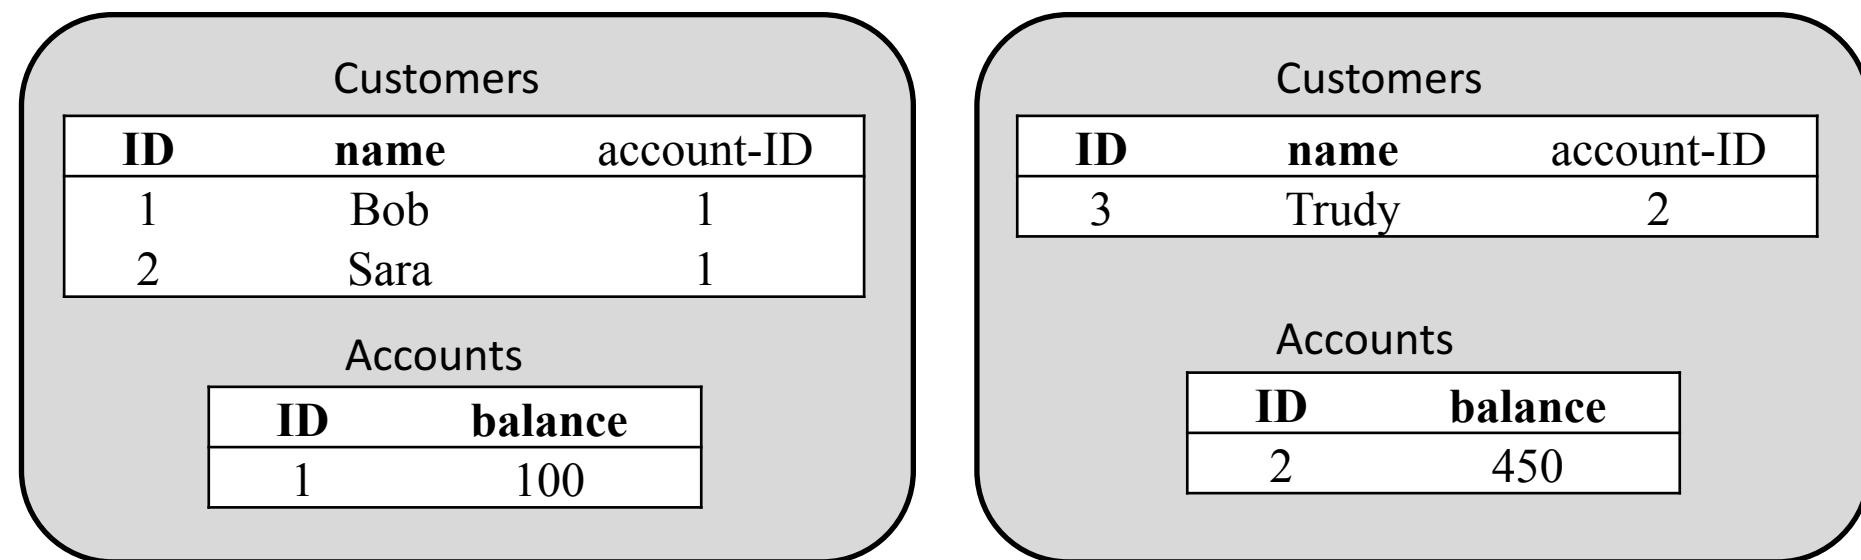
e.g. applications involving money



NewSQL

Data is normalized

Tables partitioned and replicated across machines



NewSQL

Problems

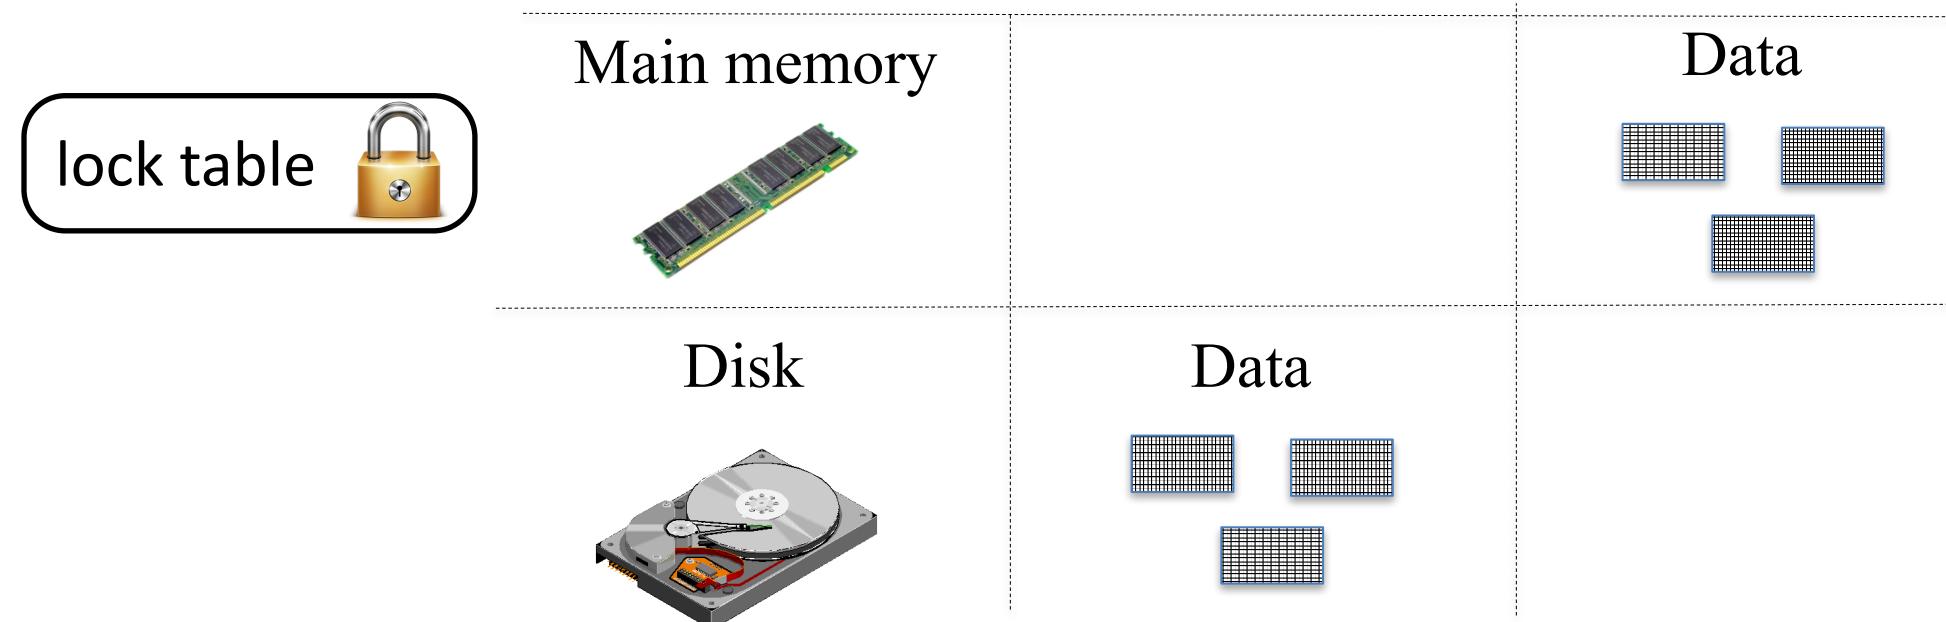
1. ACID
 2. Consistency
 3. Joins
- 
- Reengineering

Locking

concurrency introduced due to disk

Today data is in main memory

Remove concurrency and locking. Serial transactions

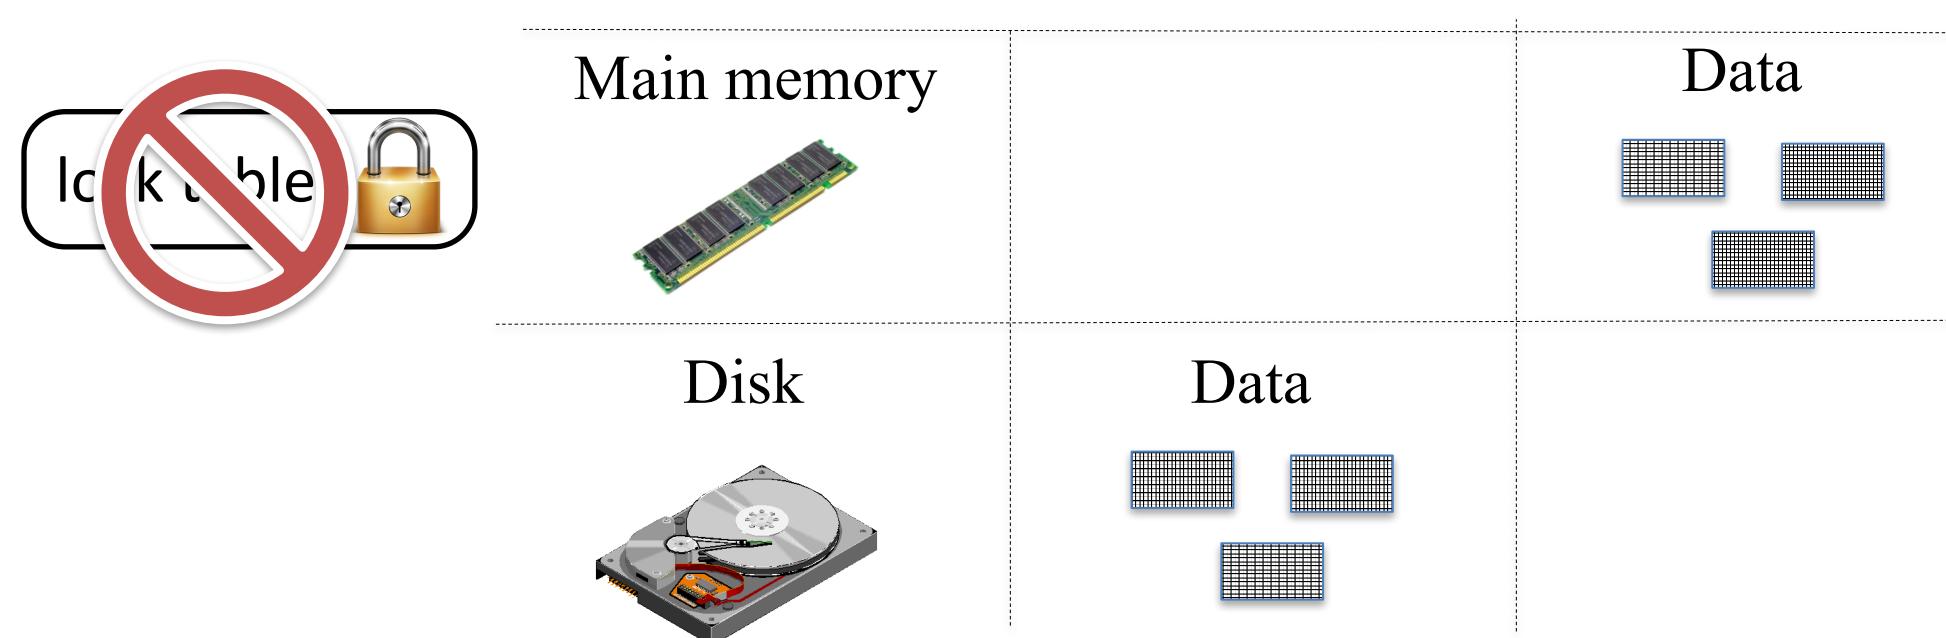


Locking

concurrency introduced due to disk

Today data is in main memory

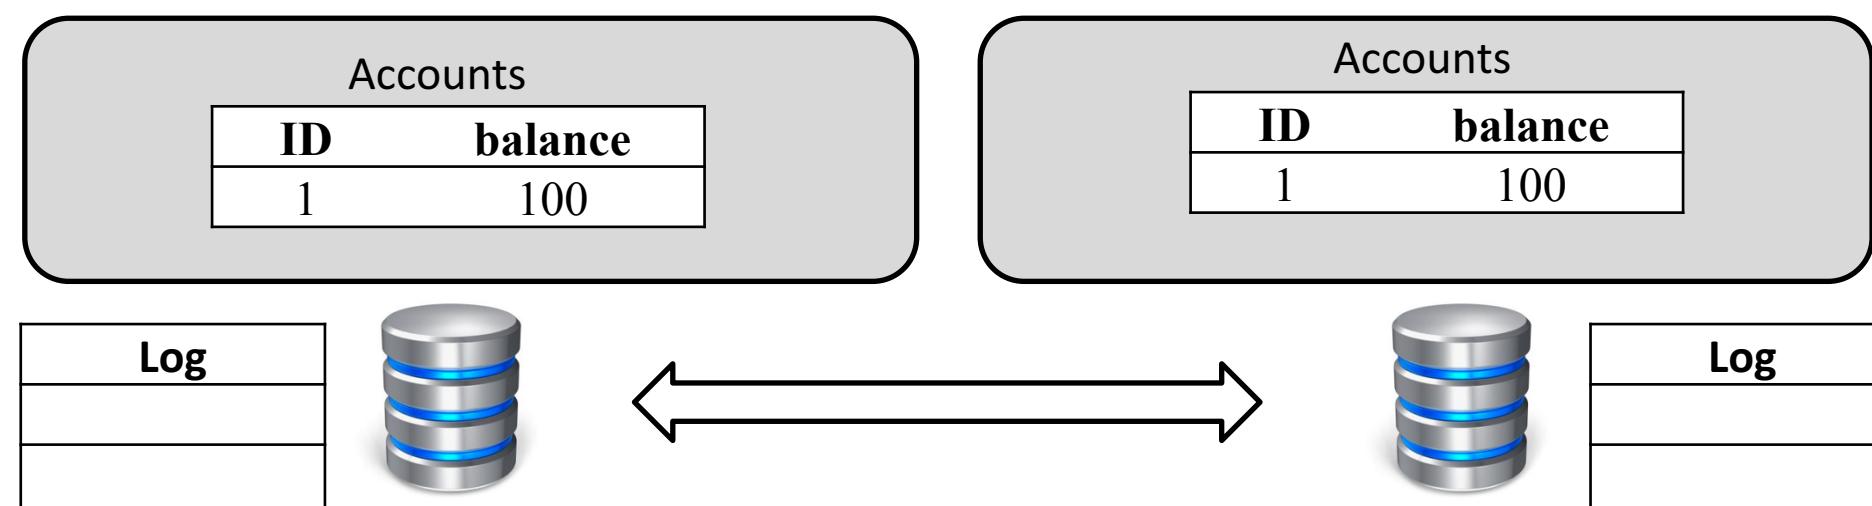
Remove concurrency and locking. Serial transactions



Logging

History: log introduced for recovery

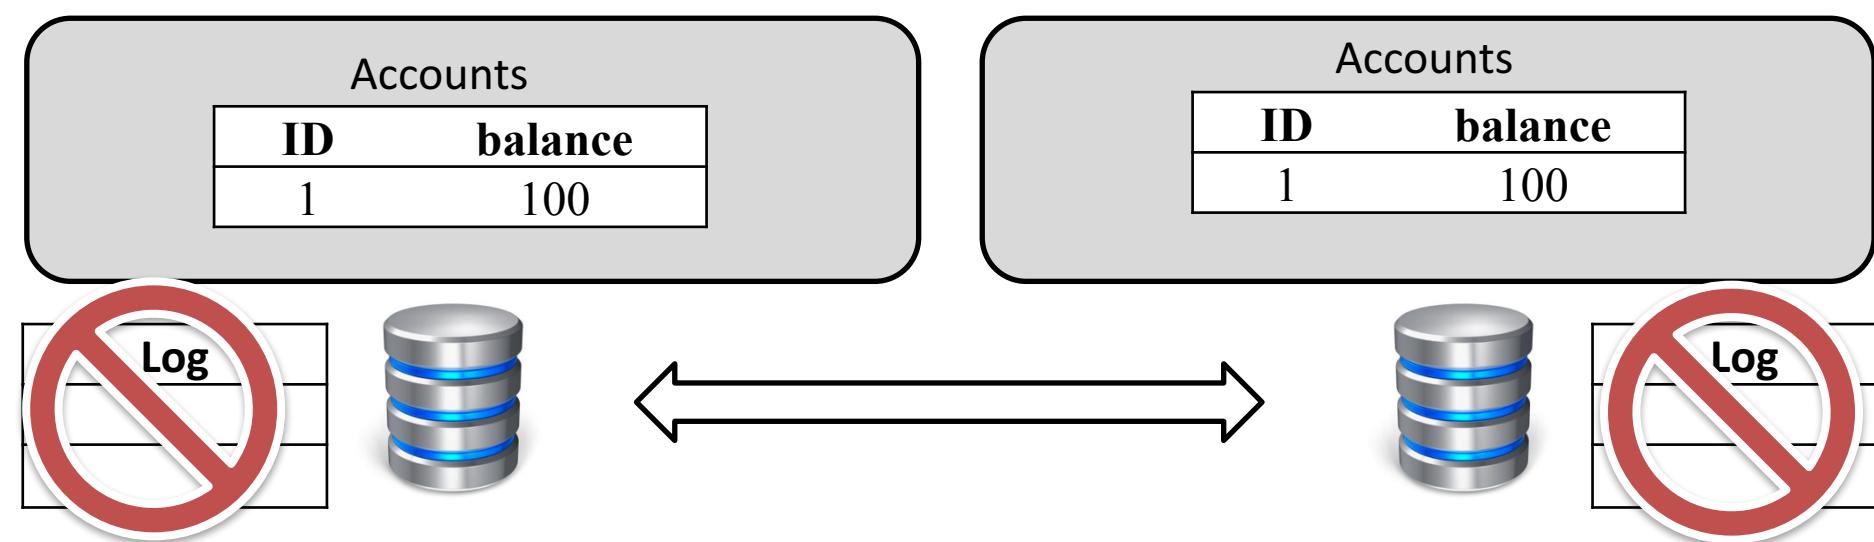
Today: replicas exist elsewhere



Logging

History: log introduced for recovery

Today: replicas exist elsewhere



NewSQL

Problems

- 1. ACID
 - 2. Consistency
 - 3. Joins
- 
- Reengineering

NewSQL

Problems

1. ACID } Reengineering
2. Consistency } Tough luck
3. Joins

NewSQL

Enforce operation order across replicas (expensive)



Deposit
Add 100

ID	balance
...	100



Interest
Multiply by 1.1

ID	balance
...	100

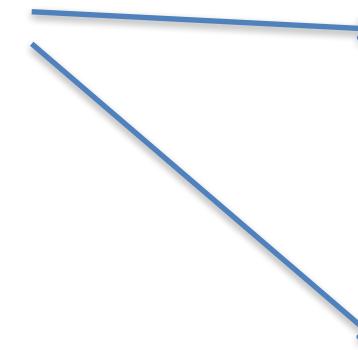


NewSQL

Enforce operation order across replicas (expensive)



Deposit
Add 100



ID	balance
...	100



Interest
Multiply by 1.1

ID	balance
...	100

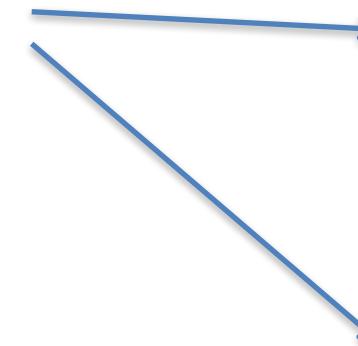


NewSQL

Enforce operation order across replicas (expensive)



Deposit
Add 100



ID	balance
...	200



Interest
Multiply by 1.1

ID	balance
...	200



NewSQL

Enforce operation order across replicas (expensive)



Deposit
Add 100

ID	balance
...	200



Interest
Multiply by 1.1

ID	balance
...	200



NewSQL

Enforce operation order across replicas (expensive)



Deposit
Add 100



Interest
Multiply by 1.1



ID	balance
...	200



ID	balance
...	200



NewSQL

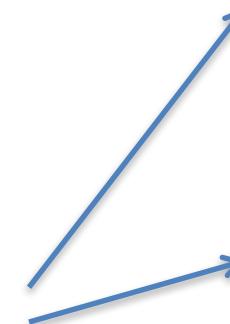
Enforce operation order across replicas (expensive)



Deposit
Add 100



Interest
Multiply by 1.1



ID	balance
...	220



ID	balance
...	220



NewSQL

Enforce operation order across replicas (expensive)



Deposit
Add 100

ID	balance
...	220



Interest
Multiply by 1.1

ID	balance
...	220



NewSQL

Problems

1. ACID } Reengineering
2. Consistency } Tough luck
3. Joins

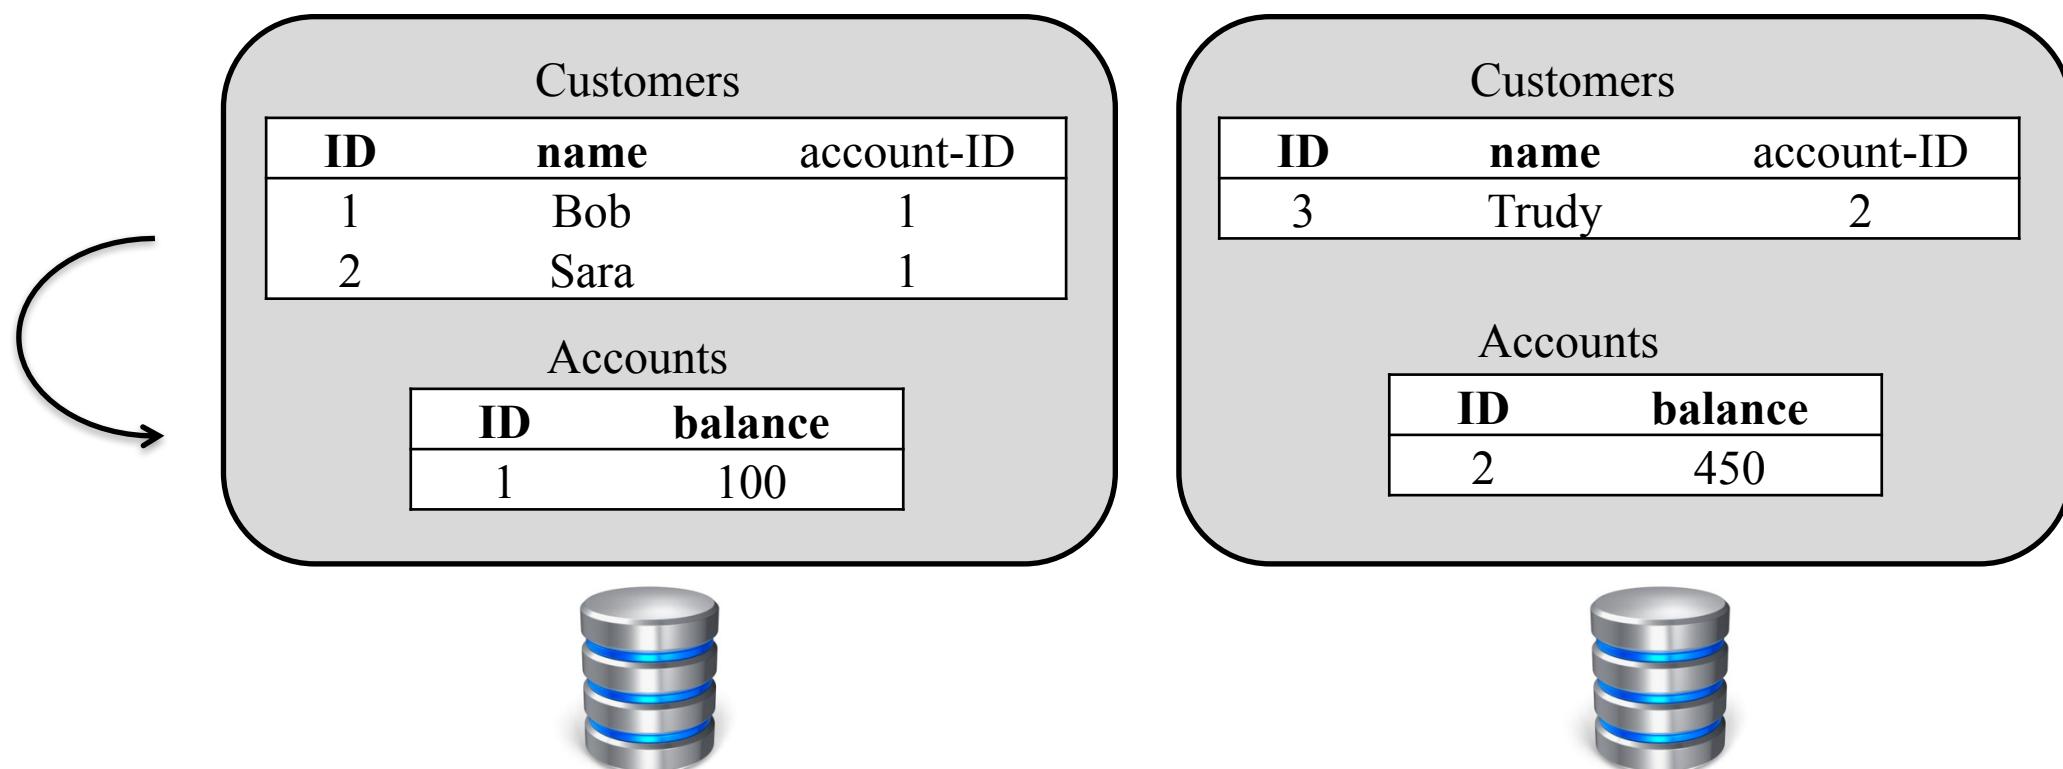
NewSQL

Problems

1. ACID } Reengineering
2. Consistency } Tough luck
3. Joins } Tough luck

NewSQL

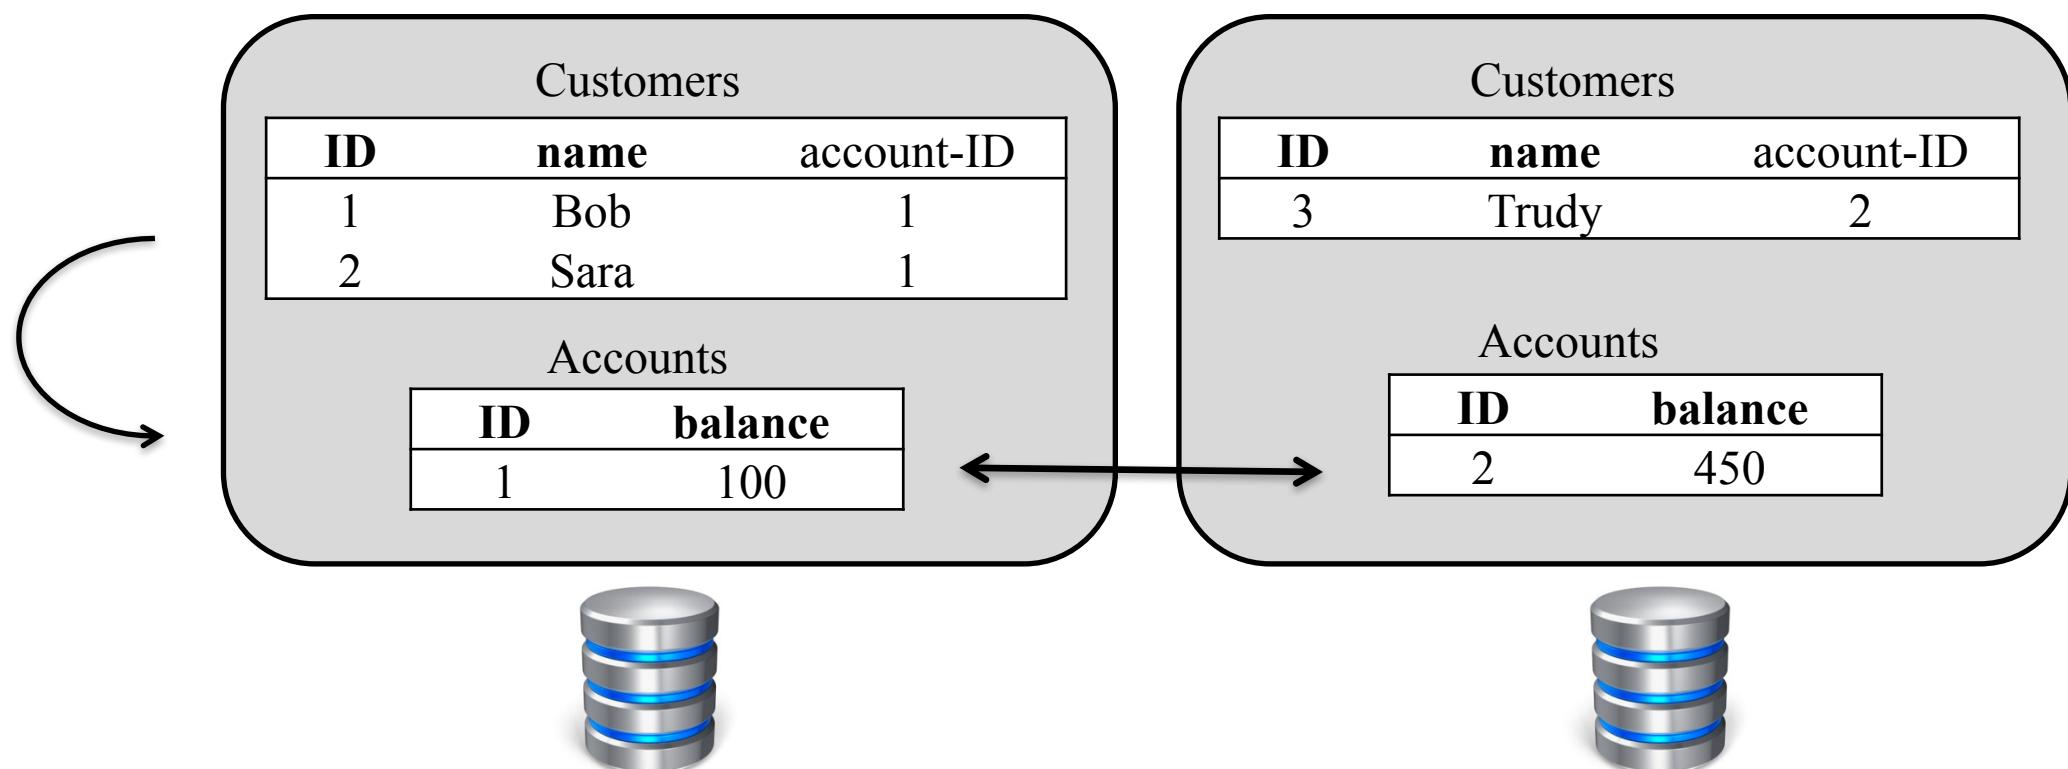
Data about same object on same machine



NewSQL

Data about same object on same machine

Transactions across machines cost (e.g. money transfers)



When to use NewSQL

Transactions span multiple data objects

Strong consistency

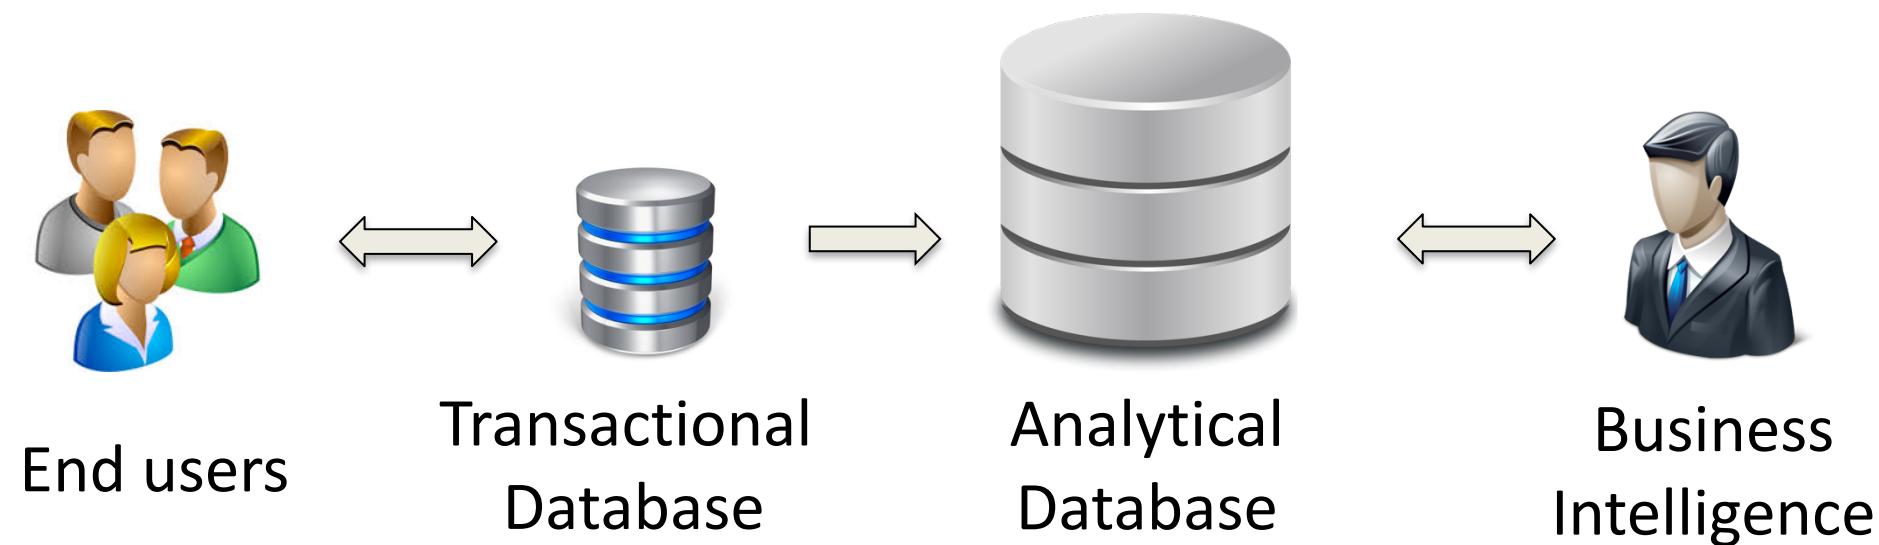
Run at scale

More expertise



Conclusion

Analytical Databases

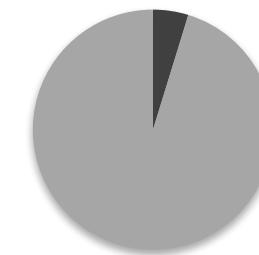


Column-Stores

Joins

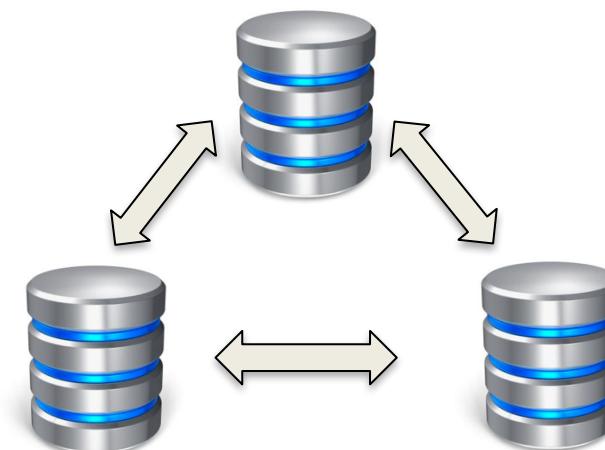
Customers			Accounts	
ID	name	acc-ID	ID	balance
2	Bob	X	X	100
...

Locking & Logging



- Useful work
- ACID

Consistency



NoSQL Databases

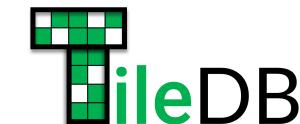
- Simplify
- De-normalize
- Weak consistency
- Non-interacting objects

NewSQL Databases

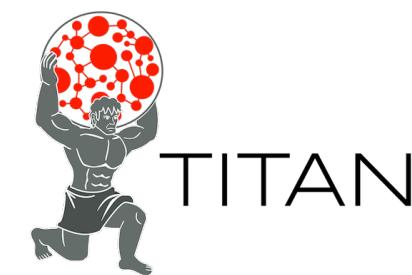
- Reengineer
- Normalize
- Strong consistency
- Interacting objects

There is more...

Scientific databases:



Graph databases



Considered NoSQL

Hopefully now you have reasoning tools

DATA SYSTEMS. LABORATORY

@ Harvard School of Engineering and Applied Sciences



Stratos Idreos



Manos
Athanassoulis



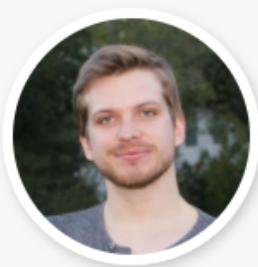
Niv Dayan



Kostas
Zoumpatianos



Michael Kester



Lukas Maas



Abdul Wasay



Brian Hentschel

prof

post-docs

phd students

<http://daslab.seas.harvard.edu/>

Thanks!!!

Questions?

Presentation: nivdayan.github.io/db_pres.pdf

MongoDB tutorial: nivdayan.github.io/mongo.pdf