

Tutorial 3


Database System Technology - Indexing

Niv Dayan - 25 Jan, 2023

Groups

Most of you are grouped up.

One new person in the course is looking for a team!



Agenda

More on SSD
write-amplification



4 exercises on
indexing



Introduce step 2 of
the project



Modeling SSD Write-Amp More Precisely

There used to be some slides here, but they are not superseded by the following more thorough and correct explanation:

https://6c7de8b5-df3d-4a03-945a-9846af236553.usrfiles.com/ugd/6c7de8_a8d4def3b5d94b4f98c2238c064f28b2.pdf

Question 1

Consider a B-tree subject to uniformly randomly distributed updates. There are 100 entries per page. The b-tree occupies 70% of the SSD, while the rest is over-provisioned. What write-amplification would you expect?

Model: $B \cdot \left(1 + \frac{1}{2} \cdot \frac{L/P}{1 - L/P}\right)$

Where L = logical data size

P = physical SSD capacity

Under what kind of workload would write-amplification for a B-tree be significantly lower?

Question 1

Consider a B-tree subject to uniformly randomly distributed updates. There are 100 entries per page. The b-tree occupies 70% of the SSD, while the rest is over-provisioned. What write-amplification would you expect?

Model:
$$B \cdot \left(1 + \frac{1}{2} \cdot \frac{L/P}{1 - L/P}\right) = 100 \cdot \left(1 + \frac{1}{2} \cdot \frac{0.7}{1 - 0.7}\right) = \mathbf{167}$$

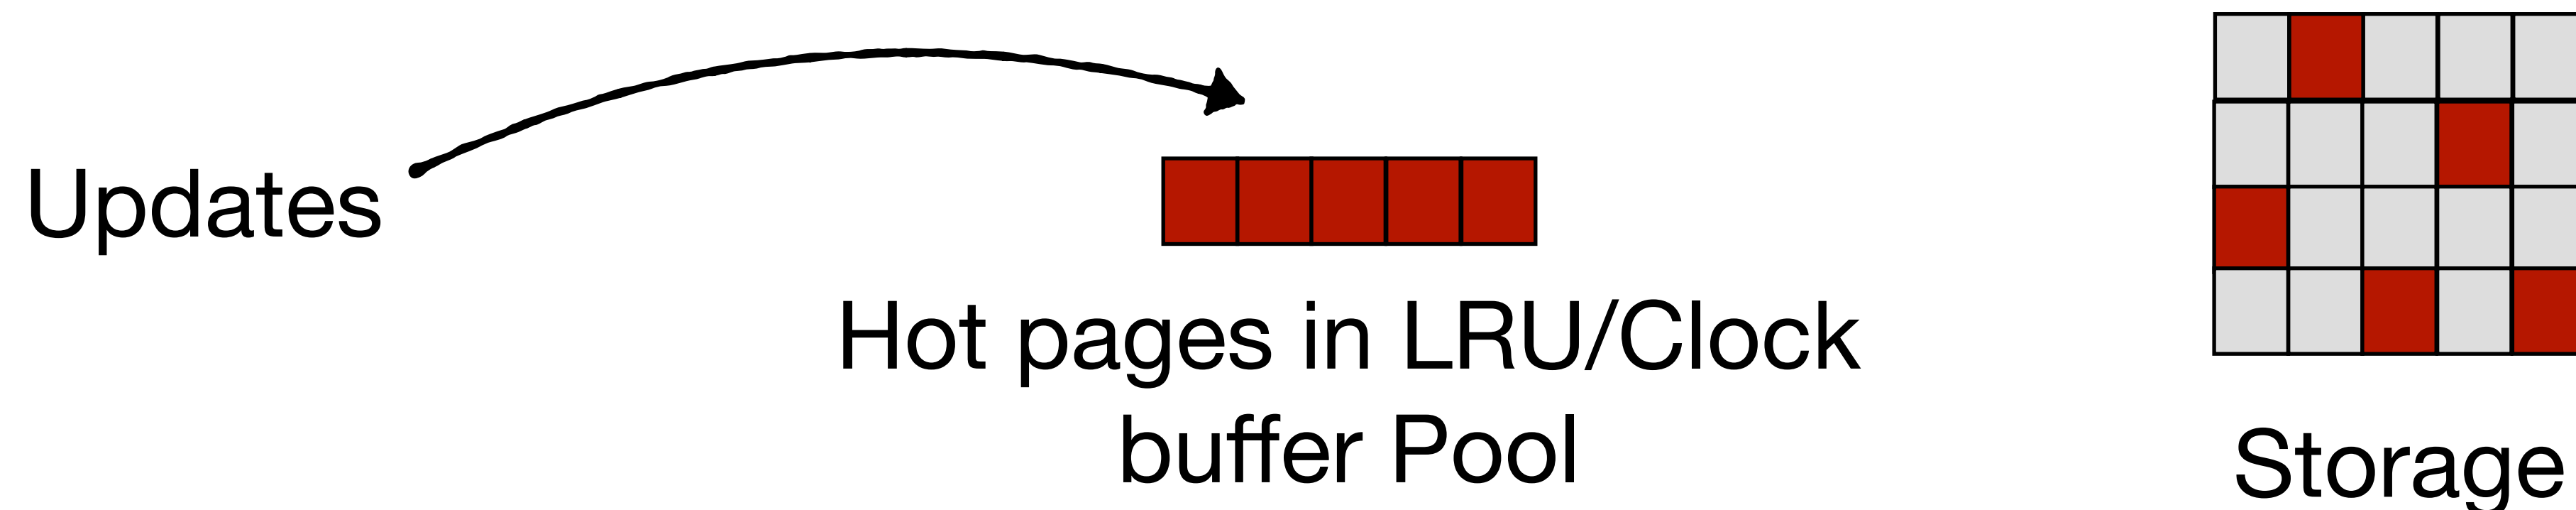
Under what kind of workload would write-amplification for a B-tree be significantly lower?

Question 1

Consider a B-tree subject to uniformly randomly distributed updates. There are 100 entries per page. The b-tree occupies 70% of the SSD, while the rest is over-provisioned. What write-amplification would you expect?

Under what kind of workload would write-amplification for a B-tree be significantly lower?

When a workload exhibits spatial and/or temporal locality: some areas of the logical address space are hot



Question 1

Consider a B-tree subject to uniformly randomly distributed updates. There are 100 entries per page. The b-tree occupies 70% of the SSD, while the rest is over-provisioned. What write-amplification would you expect?


Under what kind of workload would write-amplification for a B-tree be significantly lower?

When a workload exhibits spatial and/or temporal locality: some areas of the logical address space are hot

**This argument does not hold for extendible hashing as entries that are adjacent logically are distributed randomly in the hash table!
This is a disadvantage of hash vs tree indexes.**

Question 2

Consider the possibility of making each B-tree node take up two rather than just one flash pages (i.e., 8KB rather than 4KB). This can make the tree shallower. Is this a good idea for flash? How about on Disk?

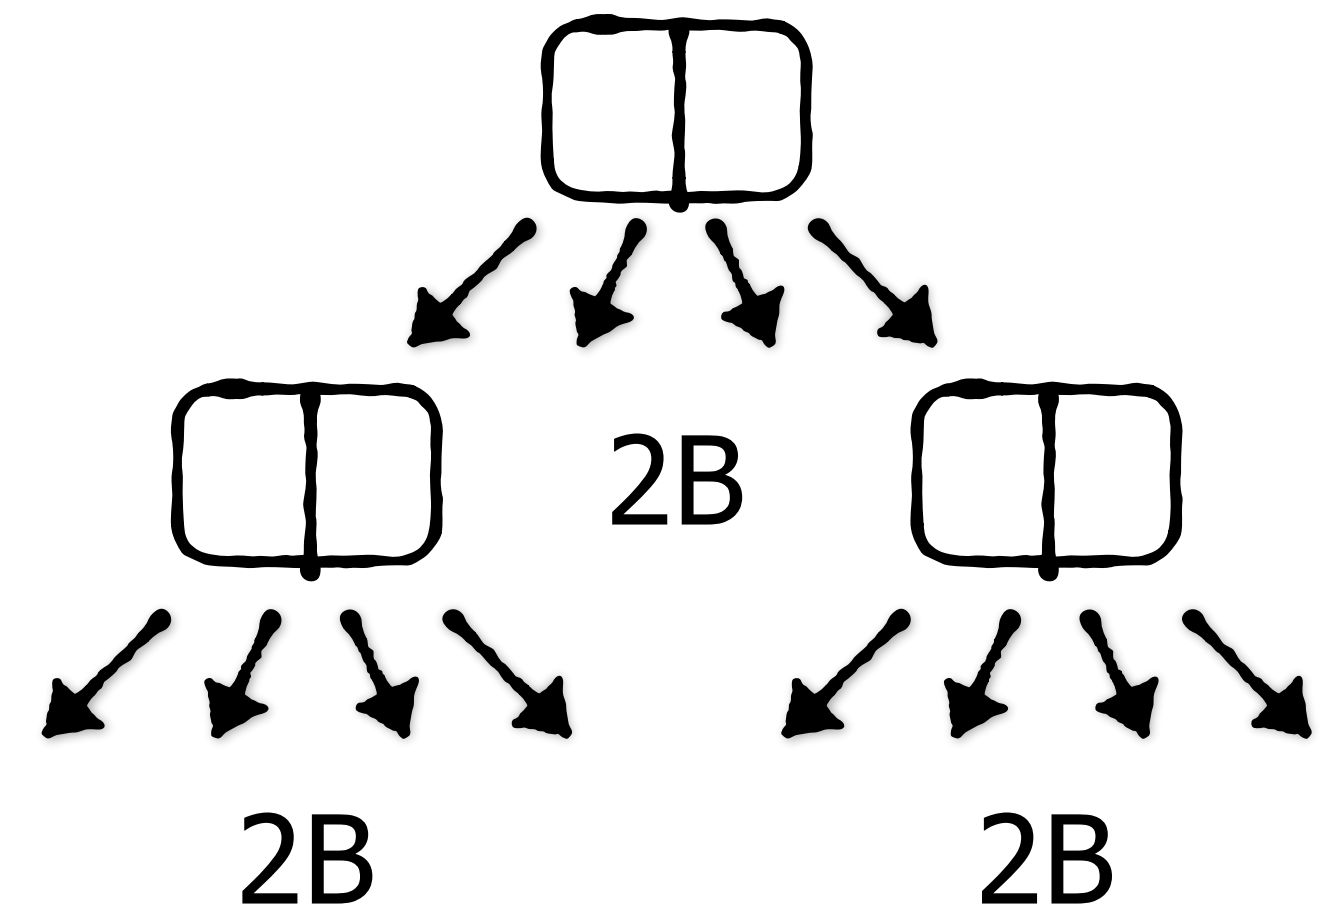


Question 2

Consider the possibility of making each B-tree node take up two rather than just one flash pages. This can make the tree shallower. Is this a good idea? How about on Disk?

On SSD, cost is measured as # pages accessed

$$= 2 * \log_{2B}(N)$$



Question 2

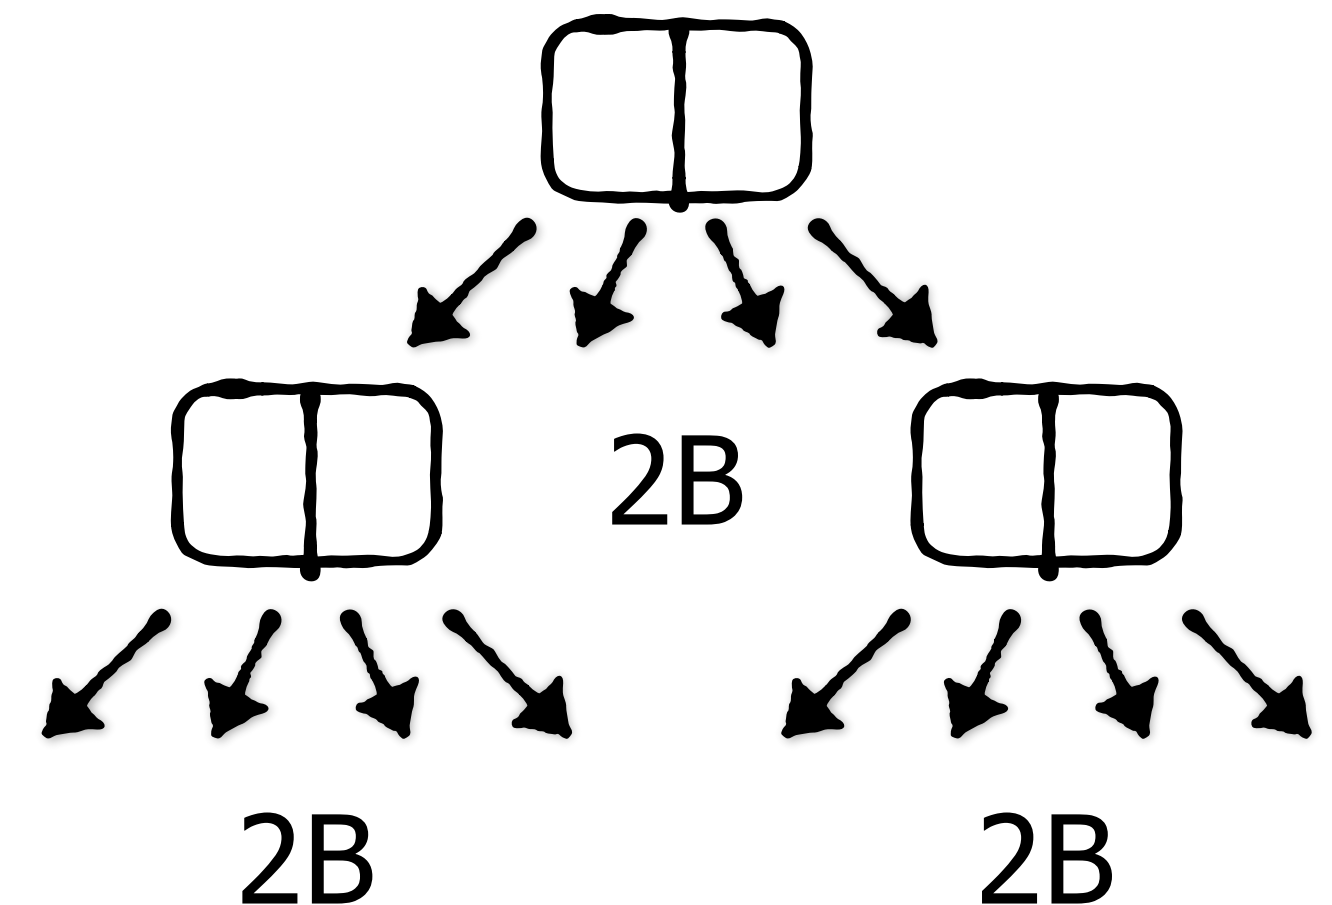
Consider the possibility of making each B-tree node take up two rather than just one flash pages. This can make the tree shallower. Is this a good idea? How about on Disk?

On SSD, cost is measured as # pages accessed

$$2 * \log_{2B}(N) \leq \log_B(N)$$



**Condition for being cheaper
than standard B-tree**




Question 2

Consider the possibility of making each B-tree node take up two rather than just one flash pages. This can make the tree shallower. Is this a good idea? How about on Disk?

On SSD, cost is measured as # pages accessed

$$2 * \log_{2B}(N) \leq \log_B(N)$$

Simplifies to: $B \leq 2$



Question 2


Consider the possibility of making each B-tree node take up two rather than just one flash pages. This can make the tree shallower. Is this a good idea? How about on Disk?

On SSD, cost is measured as # pages accessed

$$2 * \log_{2B}(N) \leq \log_B(N)$$

Simplifies to: $B \leq 2$

**But B is typically larger. So
it's not generally a good idea.**



Question 2

Consider the possibility of making each B-tree node take up two rather than just one flash pages. This can make the tree shallower. Is this a good idea? How about on Disk?

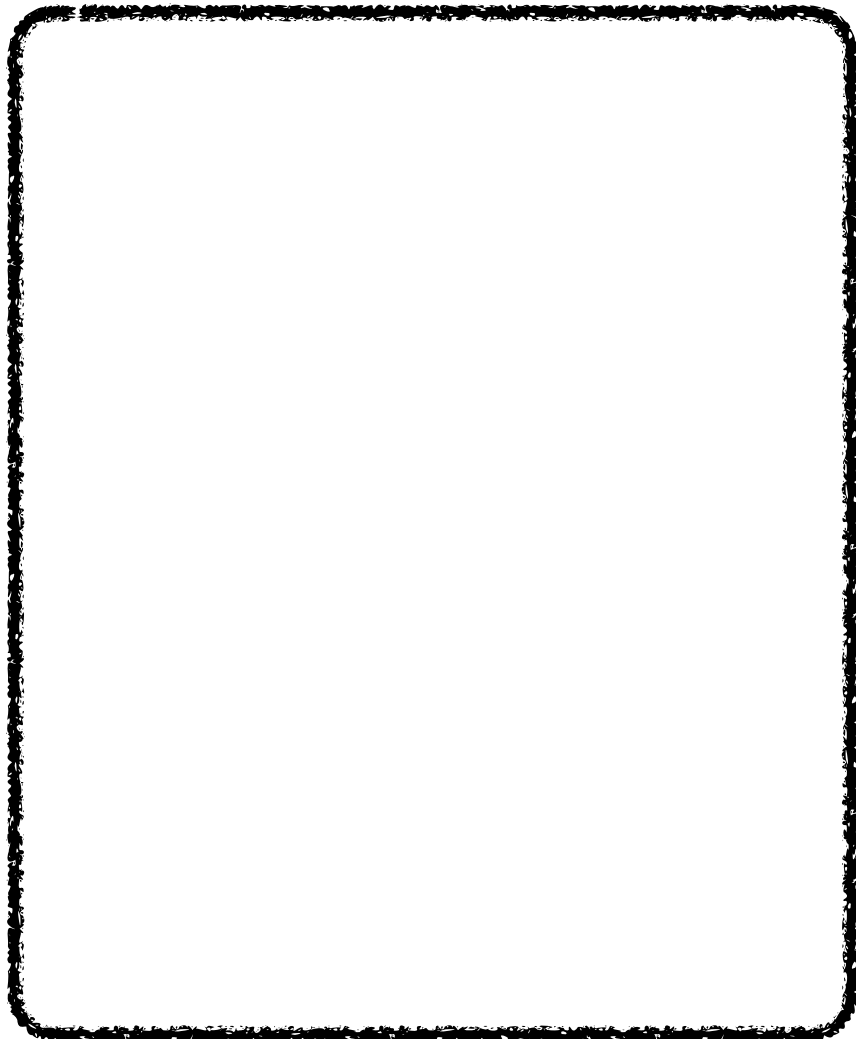


On disk, seek & rotational delay dominate, while data transfer is negligible. So this is a good idea (enlarging the node size by a multiplicative factor of B will approx. halve the depth). Likely incur diminishing returns beyond that.

Question 3

Consider a table with columns A, B and C. Suppose we employ buffered inserts at a cost of $O(1/B)$ each.

A B C



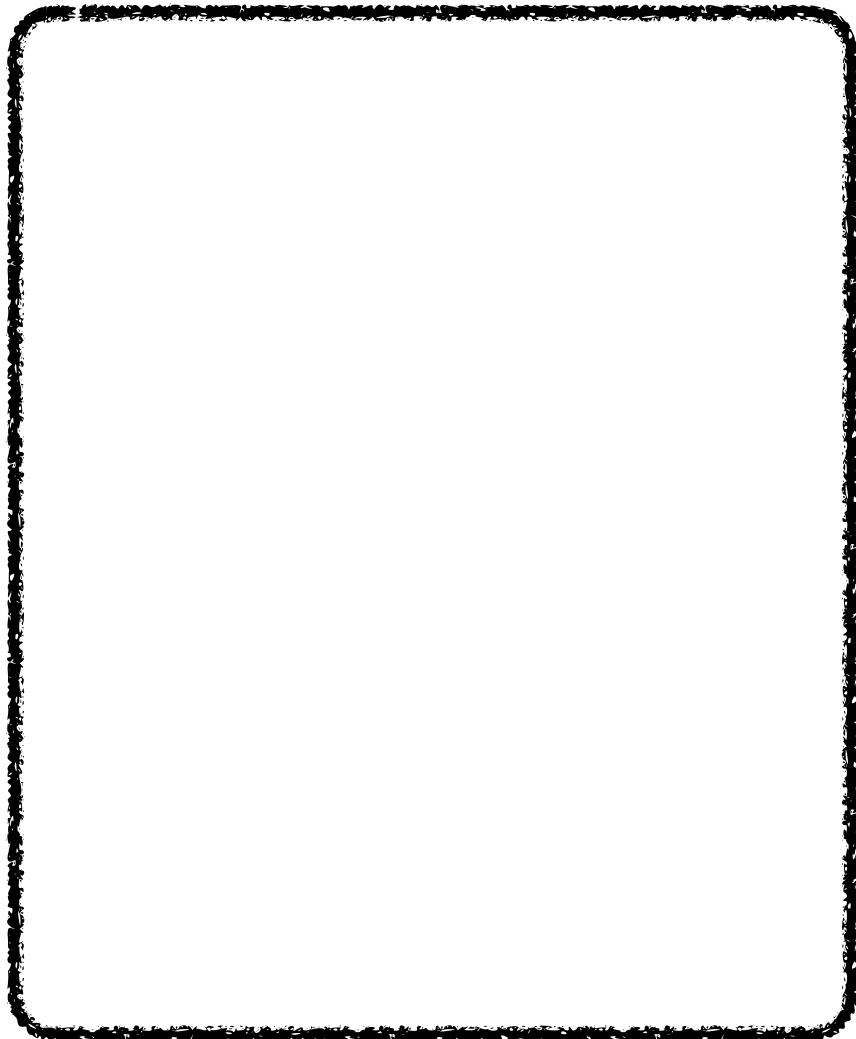
The workload consists of:

- 50% Select * from table where A = "... " Return 1 row each
- 50% Insert (, ,)

Question 3

Consider a table with columns A, B and C. Suppose we employ buffered inserts at a cost of $O(1/B)$ each.

A B C



The workload consists of:

50% Select * from table where A = "... " Return 1 row each


50% Insert (, ,)

Should we employ a B-tree index on any of the columns?
Estimate the overall I/O cost of both queries with and without it.

Question 3

Consider a table with columns A, B and C. Suppose we employ buffered inserts at a cost of $O(1/B)$ each.

A B C



50% Select * from table where A = "... " Return 1 row each

50% Insert (, ,)

Should we employ a B-tree index on any of the columns?
Estimate the overall I/O cost of both queries with and without out it.


Indexing A significantly reduces overall costs.

I/O cost without index: $0.5 * N/B + 0.5 * 1/B$ $\approx N/B$

I/O cost with index: $0.5 * \log_B N + 0.5 * \log_B N$ $\approx \log_B N$

Question 4

Consider a table with columns A, B and C.



50% Select A from table where A = "..."/> Returns 1 row each

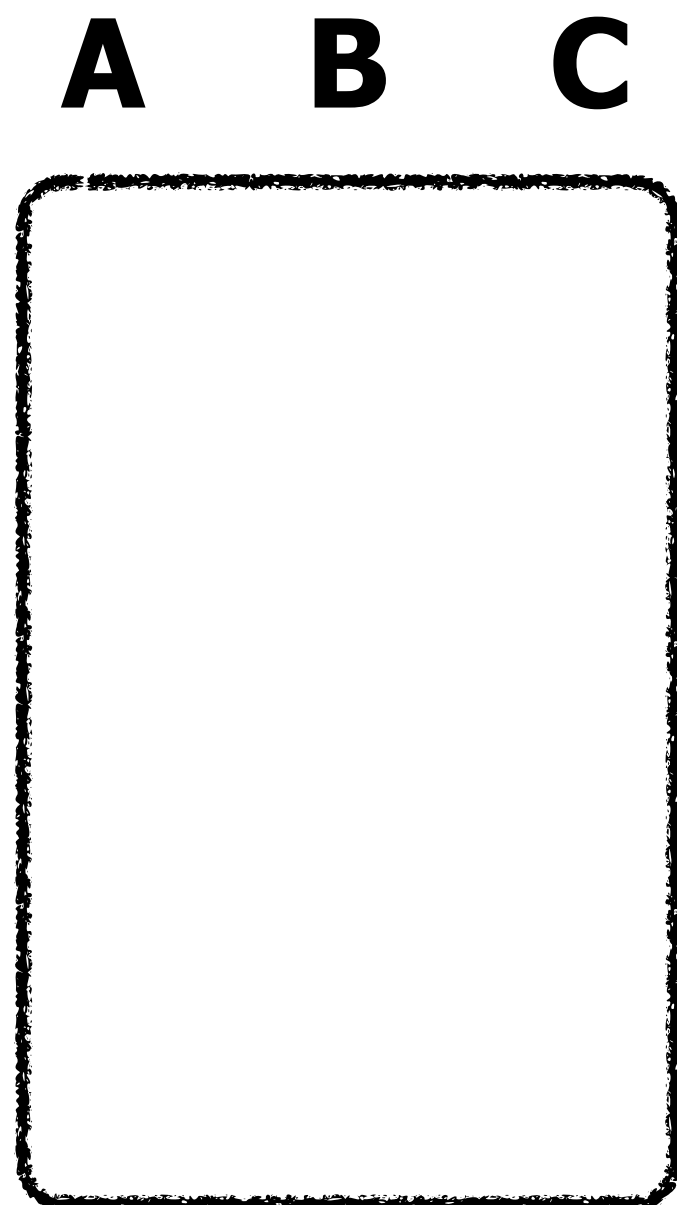
50% Select * from table where B > x and B < y Returns avg. S=10 rows

How should we index this table? B-tree or extendible hashing?

Clustered vs. unclustered? Estimate worst-case I/O cost with your plan for each query with these indexes assuming $N=2^{40}$ and $B=2^{10}$

Question 4

Consider a table with columns A, B and C.



50% Select A from table where A = "... Returns 1 row each

50% Select * from table where B > x and B < y Returns avg. S=10 rows

How should we index this table? B-tree or extendible hashing?

Clustered vs. unclustered? Estimate worst-case I/O cost with your plan for each query with these indexes assuming $N=2^{40}$ and $B=2^{10}$

Clustered B-tree on B

$$\log_2(N) + S/B$$

$$4 + 1$$

Extendible Hash table on A

1-2, assuming directory is in memory and data is evenly distributed